

CS-GY 6953 / ECE-GY 7123
Deep Learning

Homework #2

Artem Shlepchenko
as14836

Problem 1

Importing necessary library.

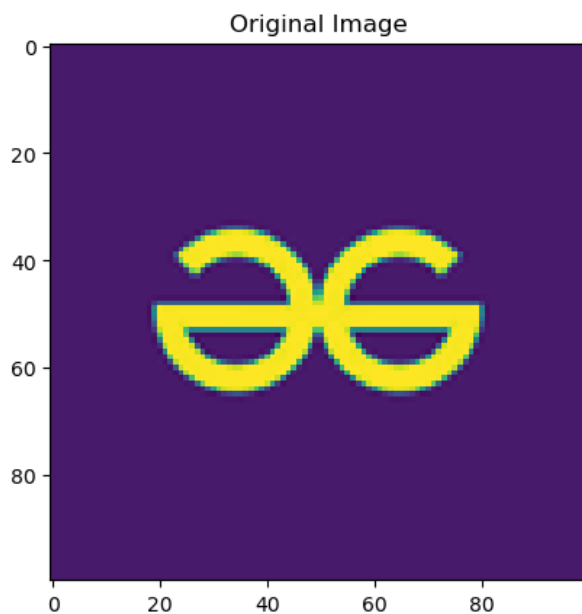
```
from PIL import Image, ImageFilter
import numpy as np
import matplotlib.pyplot as plt
import urllib.request
```

Loading an image(example).

```
urllib.request.urlretrieve(
    'https://media.geeksforgeeks.org/wp-content/uploads/20210318103632/gfg-300x300.png',
    "gfg.png")
img = np.array(Image.open("gfg.png").convert('L').resize((100, 100)))
# Print the image size
print(img.shape)
```

Displaying the image.

```
plt.figure()
plt.imshow(img)
plt.title("Original Image")
plt.show()
```



a) Write down the weights of w which acts as a blurring filter, i.e., the output is a blurry form in the input.

```
inputSize, padding, filterSize, stride = 100, 1, 3, 1

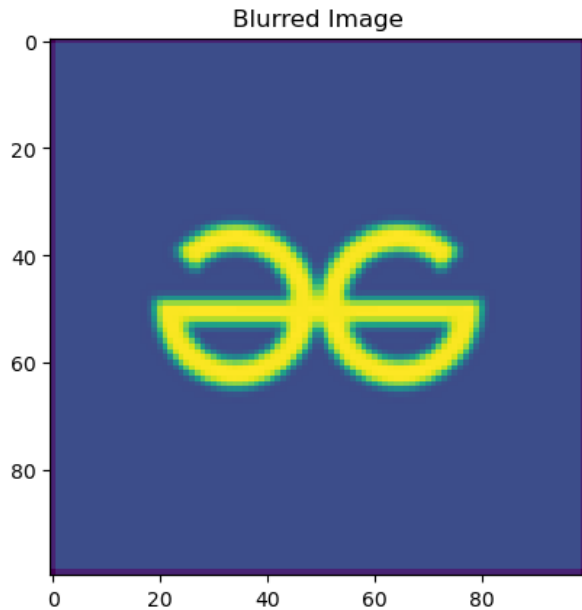
# This filter works by giving more weight to the central pixel and its immediate
# neighbors, resulting in a stronger blurring effect than the simple averaging
# of a box filter.

# The values in the filter represent the weights or coefficients that are multiplied
# with the pixel values in the corresponding positions to produce the filtered output.
# In this case, the central pixel in the filter has a weight of 4, while the surrounding
# pixels have a weight of 2. This means that the pixel values in the output image are
# more strongly influenced by the nearby pixels than those farther away, resulting in
# a more pronounced blurring effect.
filter = np.array([[1, 2, 1], [2, 4, 2], [1, 2, 1]])

#calculating size of the output image by the convolution formula
targetSize = int(np.floor(((inputSize + 2*padding - filterSize) / stride)) + 1)
output = np.zeros((targetSize, targetSize))

#Padding input image so that size of output image is same as input image.
inputImage = np.pad(img, padding, mode = 'constant')
for i in range(targetSize):
    for j in range(targetSize):
        subset = np.array(inputImage)[i:i+filterSize, j:j+filterSize]
        output[i, j] = np.sum(np.multiply(subset, filter))

plt.figure()
plt.imshow(output)
plt.title("Blurred Image")
plt.show()
```



b) Write down the weights of w which acts as a sharpening filter in the horizontal direction.

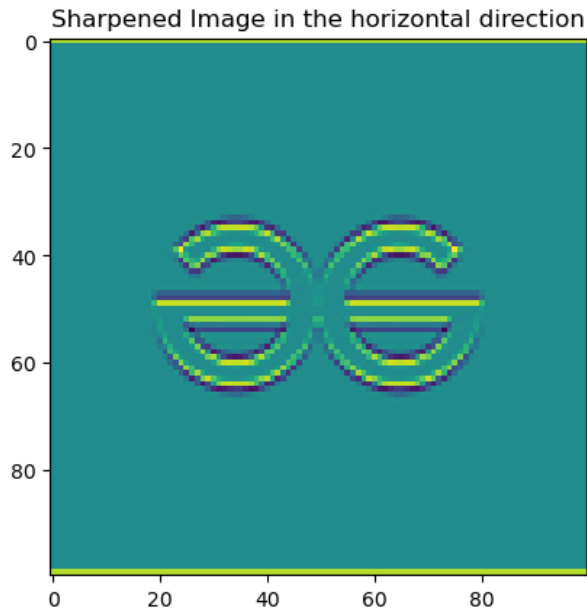
```
inputSize, padding, filterSize, stride = 100, 1, 3, 1

# This filter detects horizontal edges and enhances them, making the image appear sharper.
# The values in the filter represent the weights or coefficients that are multiplied with
# the pixel values in the corresponding positions to produce the filtered output.
# Note that the sum of the values in the filter is 0, which means that the filter is a
# high-pass filter that preserves the overall brightness of the image.
filter = np.array([[0, -1, 0], [0, 2, 0], [0, -1, 0]])

#calculating size of the output image by the convolution formula
targetSize = int(np.floor(((inputSize + 2*padding - filterSize) / stride)) + 1)
output = np.zeros((targetSize, targetSize))

#Padding input image so that size of output image is same as input image.
inputImage = np.pad(img, padding, mode = 'constant')
for i in range(targetSize):
    for j in range(targetSize):
        subset = np.array(inputImage)[i:i+filterSize, j:j+filterSize]
        output[i, j] = np.sum(np.multiply(subset, filter))

plt.figure()
plt.imshow(output)
plt.title("Sharpened Image in the horizontal direction")
plt.show()
```



c) Write down the weights of w which acts as a sharpening filter in the vertical direction.

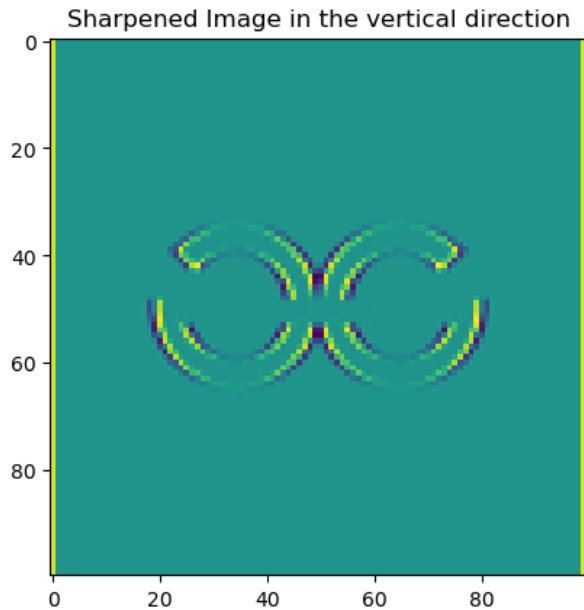
```
inputSize, padding, filterSize, stride = 100, 1, 3, 1

# This filter detects vertical edges and enhances them, making the image appear sharper.
# The values in the filter represent the weights or coefficients that are multiplied
# with the pixel values in the corresponding positions to produce the filtered output.
# Note that the sum of the values in the filter is 0, which means that the filter is
# a high-pass filter that preserves the overall brightness of the image.
filter = np.array([[0, 0, 0], [-1, 2, -1], [0, 0, 0]])

#calculating size of the output image by the convolution formula
targetSize = int(np.floor(((inputSize + 2*padding - filterSize) / stride)) + 1)
output = np.zeros((targetSize, targetSize))

#Padding input image so that size of output image is same as input image.
inputImage = np.pad(img, padding, mode = 'constant')
for i in range(targetSize):
    for j in range(targetSize):
        subset = np.array(inputImage)[i:i+filterSize, j:j+filterSize]
        output[i, j] = np.sum(np.multiply(subset, filter))

plt.figure()
plt.imshow(output)
plt.title("Sharpened Image in the vertical direction")
plt.show()
```



d) Write down the weights of w which act as a sharpening filter in a diagonal (bottom-left to top-right) direction.

```
inputSize, padding, filterSize, stride = 100, 1, 3, 1

# This filter detects diagonal edges in the bottom-left to top-right direction and
# enhances them, making the image appear sharper. The values in the filter represent
# the weights or coefficients that are multiplied with the pixel values in the
# corresponding positions to produce the filtered output.

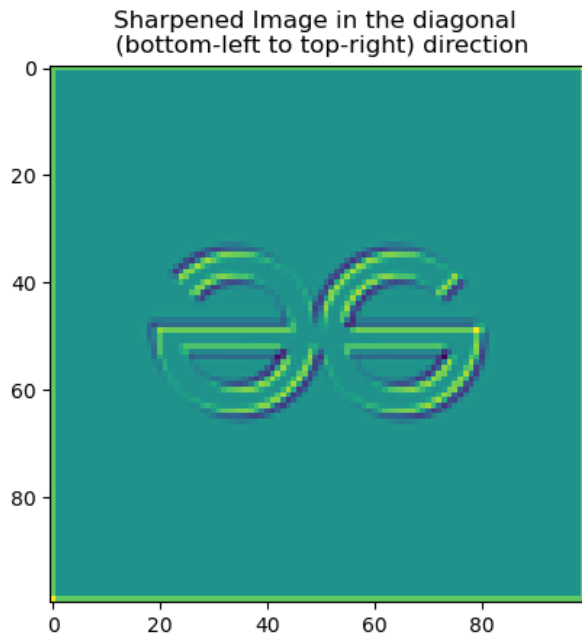
# Note that the sum of the values in the filter is 0, which means that the filter
# is a high-pass filter that preserves the overall brightness of the image.
filter = np.array([[ -1,  0,  0], [ 0,  2,  0], [ 0,  0, -1]])

#calculating size of the output image by the convolution formula
targetSize = int(np.floor(((inputSize + 2*padding - filterSize) / stride)) + 1)
output = np.zeros((targetSize, targetSize))

#Padding input image so that size of output image is same as input image.
inputImage = np.pad(img, padding, mode = 'constant')
for i in range(targetSize):
    for j in range(targetSize):
        subset = np.array(inputImage)[i:i+filterSize, j:j+filterSize]
        output[i, j] = np.sum(np.multiply(subset, filter))

plt.figure()
plt.imshow(output)
```

```
plt.title("Sharpened Image in the diagonal \n (bottom-left to top-right) direction")  
plt.show()
```



e) Give an example of an image operation which cannot be implemented using a 3x3 convolutional filter and briefly explain why.

One example of an image operation that cannot be implemented using a 3x3 convolutional filter is the dilation operation. Dilation is a morphological image processing operation that expands the boundaries of the object in an image.

In dilation, a structuring element is placed on each pixel of the image, and the pixel values are modified based on the presence of neighboring pixels within the structuring element. This operation requires a structuring element larger than 3x3 and hence cannot be implemented using a 3x3 convolutional filter.

The dilation operation can be implemented using other morphological operations, such as erosion and opening, which can be implemented using convolutional filters. But dilation operation cannot be implemented using only a 3x3 convolutional filter.

Problem 2

a)

The ℓ_2 regularised loss would be

$$L_2(w) = L(w) + \lambda(\|w\|_2)$$

b)

$$\frac{\partial L_2(w)}{\partial w} = \frac{\partial L(w)}{\partial w} + \frac{2\lambda w}{(\|w\|_2)}$$

$$W = W - \frac{\partial L_2(w)}{\partial w} = w - \left(\frac{\partial L(w)}{\partial w} + \frac{2\lambda w}{(\|w\|_2)} \right) = w \left(1 - \frac{2\lambda}{(\|w\|_2)} \right) - \frac{\partial L(w)}{\partial w}$$

c)

The expression above shows that weights are shrunk before applying the descent update.

d)

λ should be higher initially so that the weights converge faster, but once we got close enough, they should change slowly so that the weights stabilise instead of jumping and diverging.

So, we should choose λ such that it should adapt to the distance from the final solution.

Problem 3

a)

The definition IoU for any two bounding boxes is given by:


$$IoU(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

Since the right hand side is non-negative, this number has to be bigger (or equal to) 0.

Moreover, $A \cap B \subseteq A \cup B$, and hence the numerator has to be no bigger than the denominator. Therefore, the IoU metric is bounded between 0 and 1 (inclusive).

b)

Let's take two square boxes A and B with identical sizes, both aligned at the same horizontal level. Then, fix B and imagine "sliding" A from left to right. As A moves, the IoU will start from zero (no overlap), increase (until there is perfect overlap), and then decrease (until there is no overlap again). So, if we plot IoU as a function of horizontal displacement, we should get a curve like this:




which has 3 kinks and hence is non-differentiable.

Problem 4

AlexNet

In this problem, you are asked to train a deep convolutional neural network to perform image classification. In fact, this is a slight variation of a network called AlexNet. This is a landmark model in deep learning, and arguably kickstarted the current (and ongoing, and massive) wave of innovation in modern AI when its results were first presented in 2012. AlexNet was the first real-world demonstration of a deep classifier that was trained end-to-end on data and that outperformed all other ML models thus far.

We will train AlexNet using the CIFAR10 dataset, which consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. The classes are: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck.



A lot of the code you will need is already provided in this notebook; all you need to do is to fill in the missing pieces, and interpret your results.

Warning: AlexNet takes a good amount of time to train (~1 minute per epoch on Google Colab). So please budget enough time to do this homework.

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.optim.lr_scheduler import _LRScheduler
import torch.utils.data as data

import torchvision.transforms as transforms
import torchvision.datasets as datasets

from sklearn import decomposition
from sklearn import manifold
from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
import matplotlib.pyplot as plt
import numpy as np

import copy
import random
import time
```

```
SEED = 1234

random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
torch.backends.cudnn.deterministic = True
```

Loading and Preparing the Data

Our dataset is made up of color images but three color channels (red, green and blue), compared to MNIST's black and white images with a single color channel. To normalize our data we need to calculate the means and standard deviations for each of the color channels independently, and normalize them.

```
ROOT = '.data'
train_data = datasets.CIFAR10(root = ROOT,
                              train = True,
                              download = True)
```

Files already downloaded and verified

```
# Compute means and standard deviations along the R,G,B channel
```

```
means = train_data.data.mean(axis = (0,1,2)) / 255  
stds = train_data.data.std(axis = (0,1,2)) / 255
```

Next, we will do data augmentation. For each training image we will randomly rotate it (by up to 5 degrees), flip/mirror with probability 0.5, shift by +/-1 pixel. Finally we will normalize each color channel using the means/stds we calculated above.

```
train_transforms = transforms.Compose([  
    transforms.RandomRotation(5),  
    transforms.RandomHorizontalFlip(0.5),  
    transforms.RandomCrop(32, padding = 2),  
    transforms.ToTensor(),  
    transforms.Normalize(mean = means,  
                        std = stds)  
)  
  
test_transforms = transforms.Compose([  
    transforms.ToTensor(),  
    transforms.Normalize(mean = means,  
                        std = stds)  
)
```

Next, we'll load the dataset along with the transforms defined above.

We will also create a validation set with 10% of the training samples. The validation set will be used to monitor loss along different epochs, and we will pick the model along the optimization path that performed the best, and report final test accuracy numbers using this model.

```
train_data = datasets.CIFAR10(ROOT,  
                              train = True,  
                              download = True,  
                              transform = train_transforms)  
  
test_data = datasets.CIFAR10(ROOT,  
                              train = False,  
                              download = True,  
                              transform = test_transforms)
```

Files already downloaded and verified

Files already downloaded and verified

```
VALID_RATIO = 0.9

n_train_examples = int(len(train_data) * VALID_RATIO)
n_valid_examples = len(train_data) - n_train_examples

train_data, valid_data = data.random_split(train_data,
                                           [n_train_examples, n_valid_examples])
```

```
valid_data = copy.deepcopy(valid_data)
valid_data.dataset.transform = test_transforms
```

Now, we'll create a function to plot some of the images in our dataset to see what they actually look like.

Note that by default PyTorch handles images that are arranged `[channel, height, width]`, but matplotlib expects images to be `[height, width, channel]`, hence we need to permute the dimensions of our images before plotting them.

```
def plot_images(images, labels, classes, normalize = False):

    n_images = len(images)

    rows = int(np.sqrt(n_images))
    cols = int(np.sqrt(n_images))

    fig = plt.figure(figsize = (10, 10))

    for i in range(rows*cols):

        ax = fig.add_subplot(rows, cols, i+1)

        image = images[i]

        if normalize:
            image_min = image.min()
            image_max = image.max()
            image.clamp_(min = image_min, max = image_max)
            image.add_(-image_min).div_(image_max - image_min + 1e-5)

        ax.imshow(image.permute(1, 2, 0).cpu().numpy())
        ax.set_title(classes[labels[i]])
        ax.axis('off')
```

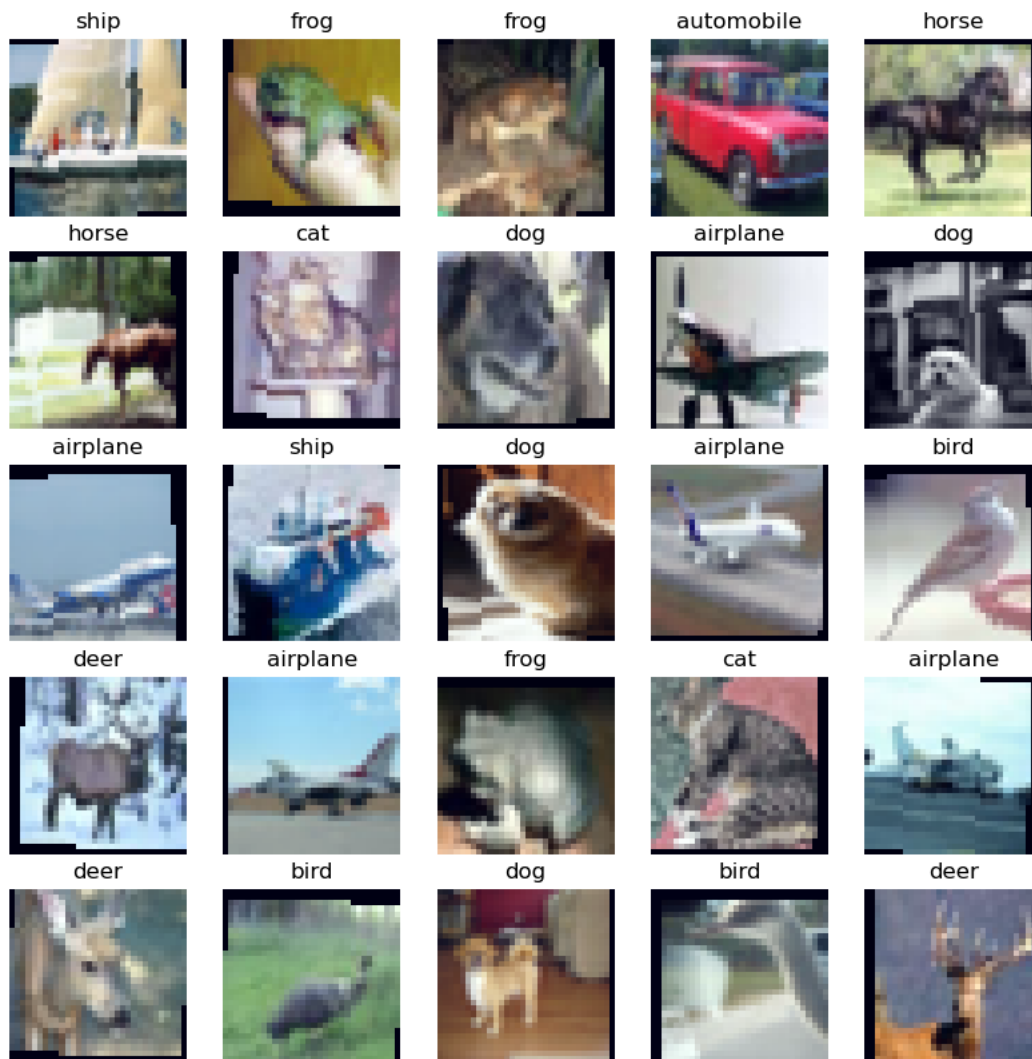
One point here: `matplotlib` is expecting the values of every pixel to be between $[0,1]$, however our normalization will cause them to be outside this range. By default `matplotlib` will then clip these values into the $[0,1]$ range. This clipping causes all of the images to look a bit weird - all of the colors are oversaturated. The solution is to normalize each image between $[0,1]$.

```
N_IMAGES = 25
```

```
images, labels = zip(*[(image, label) for image, label in
                        [train_data[i] for i in range(N_IMAGES)]])
```

```
classes = test_data.classes
```

```
plot_images(images, labels, classes, normalize = True)
```



We'll be normalizing our images by default from now on, so we'll write a function that does it for us which we can use whenever we need to renormalize an image.

```
def normalize_image(image):  
    image_min = image.min()  
    image_max = image.max()  
    image.clamp_(min = image_min, max = image_max)  
    image.add_(-image_min).div_(image_max - image_min + 1e-5)  
    return image
```

The final bit of the data processing is creating the iterators. We will use a large. Generally, a larger batch size means that our model trains faster but is a bit more susceptible to overfitting.

```
# Q1: Create data loaders for train_data, valid_data, test_data  
# Use batch size 256  
  
BATCH_SIZE = 256  
  
train_iterator = data.DataLoader(train_data, shuffle=True, batch_size=BATCH_SIZE)  
  
valid_iterator = data.DataLoader(valid_data, batch_size=BATCH_SIZE)  
  
test_iterator = data.DataLoader(test_data, batch_size=BATCH_SIZE)
```

Defining the Model

Next up is defining the model.

AlexNet will have the following architecture:

- There are 5 2D convolutional layers (which serve as feature extractors), followed by 3 linear layers (which serve as the classifier).
- All layers (except the last one) have ReLU activations. (Use inplace=True while defining your ReLUs.)
- All convolutional filter sizes have kernel size 3 x 3 and padding 1.
- Convolutional layer 1 has stride 2. All others have the default stride (1).
- Convolutional layers 1,2, and 5 are followed by a 2D maxpool of size 2.

- Linear layers 1 and 2 are preceded by Dropouts with Bernoulli parameter 0.5.
- For the convolutional layers, the number of channels is set as follows. We start with 3 channels and then proceed like this:

$3 \rightarrow 64 \rightarrow 192 \rightarrow 384 \rightarrow 256 \rightarrow 256$

In the end, if everything is correct you should get a feature map of size $2 \times 2 \times 256 = 1024$.

- the linear layers, the feature sizes are as follows:

$1024 \rightarrow 4096 \rightarrow 4096 \rightarrow 10$.

(The 10, of course, is because 10 is the number of classes in CIFAR-10).

```
class AlexNet(nn.Module):
    def __init__(self, output_dim):
        super().__init__()

        self.features = nn.Sequential(
            # Define according to the steps described above
            # Conv Layer #1
            nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3, stride=2, padding=1),
            nn.MaxPool2d(kernel_size=2),
            nn.ReLU(inplace=True),
            # Conv Layer #2
            nn.Conv2d(in_channels=64, out_channels=192, kernel_size=3, stride=1, padding=1),
            nn.MaxPool2d(kernel_size=2),
            nn.ReLU(inplace=True),
            # Conv Layer #3
            nn.Conv2d(in_channels=192, out_channels=384, kernel_size=3, stride=1, padding=1),
            nn.ReLU(inplace=True),
            # Conv Layer #4
            nn.Conv2d(in_channels=384, out_channels=256, kernel_size=3, stride=1, padding=1),
            nn.ReLU(inplace=True),
            # Conv Layer #5
            nn.Conv2d(in_channels=256, out_channels=256, kernel_size=3, stride=1, padding=1),
            nn.MaxPool2d(kernel_size=2),
            nn.ReLU(inplace=True)
        )

        self.classifier = nn.Sequential(
            # define according to the steps described above
            # Linear Layer #1
            nn.Dropout(p=0.5),
            nn.Linear(in_features=1024, out_features=4096),
            nn.ReLU(inplace=True),
```



```

        # Linear Layer #2
        nn.Dropout(p=0.5),
        nn.Linear(in_features=4096, out_features=4096),
        nn.ReLU(inplace=True),
        # Linear Layer #3
        nn.Linear(in_features=4096, out_features=output_dim)
    )

    def forward(self, x):
        x = self.features(x)
        h = x.view(x.shape[0], -1)
        x = self.classifier(h)
        return x, h

```

We'll create an instance of our model with the desired amount of classes.

```

OUTPUT_DIM = 10
model = AlexNet(OUTPUT_DIM)

```

Training the Model

We first initialize parameters in PyTorch by creating a function that takes in a PyTorch module, checking what type of module it is, and then using the `nn.init` methods to actually initialize the parameters.

For convolutional layers we will initialize using the *Kaiming Normal* scheme, also known as *He Normal*. For the linear layers we initialize using the *Xavier Normal* scheme, also known as *Glorot Normal*. For both types of layer we initialize the bias terms to zeros.

```

def initialize_parameters(m):
    if isinstance(m, nn.Conv2d):
        nn.init.kaiming_normal_(m.weight.data, nonlinearity = 'relu')
        nn.init.constant_(m.bias.data, 0)
    elif isinstance(m, nn.Linear):
        nn.init.xavier_normal_(m.weight.data, gain = nn.init.calculate_gain('relu'))
        nn.init.constant_(m.bias.data, 0)

```

We apply the initialization by using the model's `apply` method. If your definitions above are correct you should get the printed output as below.

```

model.apply(initialize_parameters)

```

```

AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
    (1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
                  ceil_mode=False)
    (2): ReLU(inplace=True)
    (3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1),
              padding=(1, 1))
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
                  ceil_mode=False)
    (5): ReLU(inplace=True)
    (6): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1),
              padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1),
              padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
              padding=(1, 1))
    (11): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
                  ceil_mode=False)
    (12): ReLU(inplace=True)
  )
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=12288, out_features=1000, bias=True)
    (2): ReLU(inplace=True)
    (3): Dropout(p=0.5, inplace=False)
    (4): Linear(in_features=1000, out_features=1000, bias=True)
    (5): ReLU(inplace=True)
    (6): Linear(in_features=1000, out_features=10, bias=True)
  )
)

```

We then define the loss function we want to use, the device we'll use and place our model and criterion on to our device.

```

optimizer = optim.Adam(model.parameters(), lr = 1e-3)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
criterion = nn.CrossEntropyLoss()

```

```
model = model.to(device)
criterion = criterion.to(device)
```

This is formatted as code

We define a function to calculate accuracy...

```
def calculate_accuracy(y_pred, y):
    top_pred = y_pred.argmax(1, keepdim = True)
    correct = top_pred.eq(y.view_as(top_pred)).sum()
    acc = correct.float() / y.shape[0]
    return acc
```

As we are using dropout we need to make sure to "turn it on" when training by using `model.train()`.

```
def train(model, iterator, optimizer, criterion, device):

    epoch_loss = 0
    epoch_acc = 0

    model.train()

    for (x, y) in iterator:

        x = x.to(device)
        y = y.to(device)

        optimizer.zero_grad()

        y_pred, _ = model(x)

        loss = criterion(y_pred, y)

        acc = calculate_accuracy(y_pred, y)

        loss.backward()

        optimizer.step()

        epoch_loss += loss.item()
        epoch_acc += acc.item()

    return epoch_loss / len(iterator), epoch_acc / len(iterator)
```

We also define an evaluation loop, making sure to "turn off" dropout with `model.eval()`.

```
def evaluate(model, iterator, criterion, device):

    epoch_loss = 0
    epoch_acc = 0

    model.eval()

    with torch.no_grad():

        for (x, y) in iterator:

            x = x.to(device)
            y = y.to(device)

            y_pred, _ = model(x)

            loss = criterion(y_pred, y)

            acc = calculate_accuracy(y_pred, y)

            epoch_loss += loss.item()
            epoch_acc += acc.item()

    return epoch_loss / len(iterator), epoch_acc / len(iterator)
```

Next, we define a function to tell us how long an epoch takes.

```
def epoch_time(start_time, end_time):
    elapsed_time = end_time - start_time
    elapsed_mins = int(elapsed_time / 60)
    elapsed_secs = int(elapsed_time - (elapsed_mins * 60))
    return elapsed_mins, elapsed_secs
```

Then, finally, we train our model.

Train it for 25 epochs (using the train dataset). At the end of each epoch, compute the validation loss and keep track of the best model. You might find the command `torch.save` helpful.

At the end you should expect to see validation losses of $\sim 76\%$ accuracy.

```
# Q3: train your model here for 25 epochs.
# Print out training and validation loss/accuracy of the model after each epoch
# Keep track of the model that achieved best validation loss thus far.
```

```

EPOCHS = 25

# Fill training code here
best_val_loss = float('inf')

for epoch in range(EPOCHS):

    start_time = time.monotonic()

    train_loss, train_acc = train(model, train_iterator, optimizer, criterion, device)
    valid_loss, valid_acc = evaluate(model, valid_iterator, criterion, device)

    if valid_loss < best_val_loss:
        best_val_loss = valid_loss
        torch.save(model.state_dict(), 'best_model.pt')

    end_time = time.monotonic()

    epoch_mins, epoch_sec = epoch_time(start_time, end_time)

    print(f'Epoch: {epoch+1:02} | Epoch Time: {epoch_mins}m {epoch_sec}s')
    print(f'\tTrain Loss: {train_loss:.3f} | Train Acc: {train_acc*100:.2f}%')
    print(f'\tVal. Loss: {valid_loss:.3f} | Val. Acc: {valid_acc*100:.2f}%')

```

```

Epoch: 01 | Epoch Time: 1m 45s
    Train Loss: 2.384 | Train Acc: 21.65%
    Val. Loss: 1.618 | Val. Acc: 38.25%
Epoch: 02 | Epoch Time: 1m 50s
    Train Loss: 1.541 | Train Acc: 43.03%
    Val. Loss: 1.374 | Val. Acc: 49.61%
Epoch: 03 | Epoch Time: 1m 51s
    Train Loss: 1.353 | Train Acc: 50.88%
    Val. Loss: 1.197 | Val. Acc: 56.91%
Epoch: 04 | Epoch Time: 1m 49s
    Train Loss: 1.263 | Train Acc: 54.77%
    Val. Loss: 1.144 | Val. Acc: 59.15%
Epoch: 05 | Epoch Time: 1m 51s
    Train Loss: 1.180 | Train Acc: 58.10%
    Val. Loss: 1.094 | Val. Acc: 60.06%
Epoch: 06 | Epoch Time: 1m 53s
    Train Loss: 1.117 | Train Acc: 60.31%
    Val. Loss: 1.059 | Val. Acc: 62.04%

```

Epoch: 07 | Epoch Time: 1m 52s
Train Loss: 1.073 | Train Acc: 62.12%
Val. Loss: 1.050 | Val. Acc: 64.64%

Epoch: 08 | Epoch Time: 1m 56s
Train Loss: 1.021 | Train Acc: 64.08%
Val. Loss: 0.970 | Val. Acc: 65.37%

Epoch: 09 | Epoch Time: 1m 57s
Train Loss: 0.978 | Train Acc: 65.95%
Val. Loss: 0.908 | Val. Acc: 67.94%

Epoch: 10 | Epoch Time: 2m 0s
Train Loss: 0.938 | Train Acc: 67.19%
Val. Loss: 0.918 | Val. Acc: 68.45%

Epoch: 11 | Epoch Time: 1m 59s
Train Loss: 0.900 | Train Acc: 68.30%
Val. Loss: 0.882 | Val. Acc: 69.49%

Epoch: 12 | Epoch Time: 1m 58s
Train Loss: 0.890 | Train Acc: 68.96%
Val. Loss: 0.868 | Val. Acc: 70.55%

Epoch: 13 | Epoch Time: 2m 2s
Train Loss: 0.847 | Train Acc: 70.55%
Val. Loss: 0.866 | Val. Acc: 70.22%

Epoch: 14 | Epoch Time: 2m 0s
Train Loss: 0.826 | Train Acc: 71.15%
Val. Loss: 0.816 | Val. Acc: 71.85%

Epoch: 15 | Epoch Time: 1m 59s
Train Loss: 0.802 | Train Acc: 71.91%
Val. Loss: 0.801 | Val. Acc: 72.59%

Epoch: 16 | Epoch Time: 2m 0s
Train Loss: 0.774 | Train Acc: 73.09%
Val. Loss: 0.819 | Val. Acc: 71.74%

Epoch: 17 | Epoch Time: 2m 2s
Train Loss: 0.765 | Train Acc: 73.60%
Val. Loss: 0.791 | Val. Acc: 73.27%

Epoch: 18 | Epoch Time: 2m 2s
Train Loss: 0.756 | Train Acc: 74.04%
Val. Loss: 0.782 | Val. Acc: 73.12%

Epoch: 19 | Epoch Time: 2m 3s
Train Loss: 0.726 | Train Acc: 74.90%
Val. Loss: 0.778 | Val. Acc: 73.49%

Epoch: 20 | Epoch Time: 2m 2s
Train Loss: 0.710 | Train Acc: 75.32%

```
Val. Loss: 0.737 | Val. Acc: 74.64%
Epoch: 21 | Epoch Time: 2m 3s
Train Loss: 0.693 | Train Acc: 76.35%
Val. Loss: 0.750 | Val. Acc: 75.41%
Epoch: 22 | Epoch Time: 2m 4s
Train Loss: 0.683 | Train Acc: 76.37%
Val. Loss: 0.736 | Val. Acc: 75.45%
Epoch: 23 | Epoch Time: 2m 1s
Train Loss: 0.674 | Train Acc: 76.82%
Val. Loss: 0.726 | Val. Acc: 75.70%
Epoch: 24 | Epoch Time: 2m 0s
Train Loss: 0.658 | Train Acc: 77.40%
Val. Loss: 0.752 | Val. Acc: 74.87%
Epoch: 25 | Epoch Time: 2m 3s
Train Loss: 0.651 | Train Acc: 77.71%
Val. Loss: 0.779 | Val. Acc: 73.01%
```

Evaluating the model

We then load the parameters of our model that achieved the best validation loss. You should expect to see ~75% accuracy of this model on the test dataset.

Finally, plot the confusion matrix of this model and comment on any interesting patterns you can observe there. For example, which two classes are confused the most?

```
# Q4: Load the best performing model, evaluate it on the test dataset, and print test
# accuracy.

model.load_state_dict(torch.load('best_model.pt'))

test_loss, test_acc = evaluate(model, test_iterator, criterion, device)

print(f'Test Loss: {test_loss:.3f} | Test Acc: {test_acc*100:.2f}%')
# Also, print out the confusion matrix.
```

```
Test Loss: 0.719 | Test Acc: 75.04%
```

```
def get_predictions(model, iterator, device):

    model.eval()

    labels = []
    probs = []

    # Q4: Fill code here.
    with torch.no_grad():

        for (x, y) in iterator:
            x = x.to(device)
            y_pred, _ = model(x)
            y_prob = F.softmax(y_pred, dim=1)

            labels.append(y.cpu())
            probs.append(y_prob.cpu())

    labels = torch.cat(labels, dim = 0)
    probs = torch.cat(probs, dim = 0)

    return labels, probs
```


```
labels, probs = get_predictions(model, test_iterator, device)
```

```
pred_labels = torch.argmax(probs, 1)
```

```
def plot_confusion_matrix(labels, pred_labels, classes):

    fig = plt.figure(figsize = (10, 10));
    ax = fig.add_subplot(1, 1, 1);
    cm = confusion_matrix(labels, pred_labels);
    cm = ConfusionMatrixDisplay(cm, display_labels = classes);
    cm.plot(values_format = 'd', cmap = 'Blues', ax = ax)
    plt.xticks(rotation = 20)
```

```
plot_confusion_matrix(labels, pred_labels, classes)
```

From the figure above we can see that the most confused classes are cat and dog.

Conclusion

That's it! As a side project (this is not for credit and won't be graded), feel free to play around with different design choices that you made while building this network.

- Whether or not to normalize the color channels in the input.
- The learning rate parameter in Adam.
- The batch size.
- The number of training epochs.
- (and if you are feeling brave – the AlexNet architecture itself.)

Problem 5

A lot of code is taken and modified from (https://pytorch.org/tutorials/intermediate/torchvision_tutorial.html) as asked in the question.

First, we need to install `pycocotools`. This library will be used for computing the evaluation metrics following the COCO metric for intersection over union.

```
%%shell

pip install cython
# Install pycocotools, the version by default in Colab
# has a bug fixed in https://github.com/cocodataset/cocoapi/pull/354
pip install -U 'git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI'
```

Let's download and extract the data, present in a zip file at https://www.cis.upenn.edu/~jshih/ped_html/PennFudanPed.zip.

```
%%shell

# download the Penn-Fudan dataset
wget https://www.cis.upenn.edu/~jshih/ped_html/PennFudanPed.zip .
# extract it in the current folder
unzip PennFudanPed.zip
```

Defining the Dataset

The dataset should inherit from the standard `torch.utils.data.Dataset` class, and implement `__len__` and `__getitem__`.

The only specificity that we require is that the dataset `__getitem__` should return:

- image: a PIL Image of size (H, W)
- target: a dict containing the following fields
 - `boxes` (`FloatTensor[N, 4]`): the coordinates of the `N` bounding boxes in `[x0, y0, x1, y1]` format, ranging from `0` to `W` and `0` to `H`
 - `labels` (`Int64Tensor[N]`): the label for each bounding box
 - `image_id` (`Int64Tensor[1]`): an image identifier. It should be unique between all the images in the dataset, and is used during evaluation

- `area` (`Tensor[N]`): The area of the bounding box. This is used during evaluation with the COCO metric, to separate the metric scores between small, medium and large boxes.
- `iscrowd` (`UInt8Tensor[N]`): instances with `iscrowd=True` will be ignored during evaluation.
- (optionally) `masks` (`UInt8Tensor[N, H, W]`): The segmentation masks for each one of the objects.
- (optionally) `keypoints` (`FloatTensor[N, K, 3]`): For each one of the `N` objects, it contains the `K` keypoints in `[x, y, visibility]` format, defining the object. `visibility=0` means that the keypoint is not visible. Note that for data augmentation, the notion of flipping a keypoint is dependent on the data representation, and you should probably adapt `references/detection/transforms.py` for your new keypoint representation

Writing a custom dataset for Penn-Fudan

Let's write a dataset for the Penn-Fudan dataset.

```
import os
import numpy as np
import torch
import torch.utils.data
from PIL import Image

class PennFudanDataset(torch.utils.data.Dataset):
    def __init__(self, root, transforms=None):
        self.root = root
        self.transforms = transforms
        # load all image files, sorting them to
        # ensure that they are aligned
        self.imgs = list(sorted(os.listdir(os.path.join(root, "PNGImages"))))
        self.masks = list(sorted(os.listdir(os.path.join(root, "PedMasks"))))

    def __getitem__(self, idx):
        # load images and masks
        img_path = os.path.join(self.root, "PNGImages", self.imgs[idx])
        mask_path = os.path.join(self.root, "PedMasks", self.masks[idx])
        img = Image.open(img_path).convert("RGB")
        # note that we haven't converted the mask to RGB,
        # because each color corresponds to a different instance
        # with 0 being background
        mask = Image.open(mask_path)

        mask = np.array(mask)
        # instances are encoded as different colors
```

```

obj_ids = np.unique(mask)
# first id is the background, so remove it
obj_ids = obj_ids[1:]

# split the color-encoded mask into a set
# of binary masks
masks = mask == obj_ids[:, None, None]

# get bounding box coordinates for each mask
num_objs = len(obj_ids)
boxes = []
for i in range(num_objs):
    pos = np.where(masks[i])
    xmin = np.min(pos[1])
    xmax = np.max(pos[1])
    ymin = np.min(pos[0])
    ymax = np.max(pos[0])
    boxes.append([xmin, ymin, xmax, ymax])

boxes = torch.as_tensor(boxes, dtype=torch.float32)
# there is only one class
labels = torch.ones((num_objs,), dtype=torch.int64)
masks = torch.as_tensor(masks, dtype=torch.uint8)

image_id = torch.tensor([idx])
area = (boxes[:, 3] - boxes[:, 1]) * (boxes[:, 2] - boxes[:, 0])
# suppose all instances are not crowd
iscrowd = torch.zeros((num_objs,), dtype=torch.int64)

target = {}
target["boxes"] = boxes
target["labels"] = labels
target["masks"] = masks
target["image_id"] = image_id
target["area"] = area
target["iscrowd"] = iscrowd

if self.transforms is not None:
    img, target = self.transforms(img, target)

return img, target

def __len__(self):
    return len(self.imgs)

```

Instance Segmentation Model with ResNet Backbone (Option 1)

```
import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor
import torchvision
from torchvision.models.detection import FasterRCNN, MaskRCNN
from torchvision.models.detection.rpn import AnchorGenerator
import torch.nn as nn

def get_instance_segmentation_model_1(num_classes):
    # load an instance segmentation model pre-trained on COCO
    model = torchvision.models.detection.maskrcnn_resnet50_fpn(pretrained=True)

    # get the number of input features for the classifier
    in_features = model.roi_heads.box_predictor.cls_score.in_features
    # replace the pre-trained head with a new one
    model.roi_heads.box_predictor = FastRCNNPredictor(in_features, num_classes)

    # now get the number of input features for the mask classifier
    in_features_mask = model.roi_heads.mask_predictor.conv5_mask.in_channels
    hidden_layer = 256
    # and replace the mask predictor with a new one
    model.roi_heads.mask_predictor = MaskRCNNPredictor(in_features_mask,
                                                         hidden_layer,
                                                         num_classes)

    return model
```

Instance Segmentation Model with MobileNet Backbone (Option 2)

```
def get_instance_segmentation_model_2(num_classes):

    backbone = torchvision.models.mobilenet_v2(pretrained=True).features
    # FasterRCNN needs to know the number of
    # output channels in a backbone. For mobilenet_v2, it's 1280
    # so we need to add it here
    backbone.out_channels = 1280

    # let's make the RPN generate 5 x 3 anchors per spatial
    # location, with 5 different sizes and 3 different aspect
    # ratios. We have a Tuple[Tuple[int]] because each feature
    # map could potentially have different sizes and
```

```

# aspect ratios
anchor_generator = AnchorGenerator(sizes=((32, 64, 128, 256, 512),),
                                   aspect_ratios=((0.5, 1.0, 2.0),))

# let's define what are the feature maps that we will
# use to perform the region of interest cropping, as well as
# the size of the crop after rescaling.
# if your backbone returns a Tensor, featmap_names is expected to
# be [0]. More generally, the backbone should return an
# OrderedDict[Tensor], and in featmap_names you can choose which
# feature maps to use.
roi_pooler = torchvision.ops.MultiScaleRoIAlign(featmap_names=['0'],
                                                output_size=7,
                                                sampling_ratio=2)

mask_roi_pooler = torchvision.ops.MultiScaleRoIAlign(
    featmap_names=['0'], output_size=14, sampling_ratio=2)

# put the pieces together inside a FasterRCNN model
model = MaskRCNN(backbone,
                  num_classes=num_classes,
                  rpn_anchor_generator=anchor_generator,
                  box_roi_pool=roi_pooler,
                  mask_roi_pooler=mask_roi_pooler)

return model

```

Training and evaluation functions

In `references/detection/`, we have a number of helper functions to simplify training and evaluating detection models. Here, we will use `references/detection/engine.py`, `references/detection/utils.py` and `references/detection/transforms.py`.

Let's copy those files (and their dependencies) in here so that they are available in the notebook

```

%%shell

# Download TorchVision repo to use some files from
# references/detection
git clone https://github.com/pytorch/vision.git
cd vision
git checkout v0.8.2

```

```
cp references/detection/utils.py ../
cp references/detection/transforms.py ../
cp references/detection/coco_eval.py ../
cp references/detection/engine.py ../
cp references/detection/coco_utils.py ../
```

Let's write some helper functions for data augmentation / transformation, which leverages the functions in `references/detection` that we have just copied:

```
from engine import train_one_epoch, evaluate
import utils
import transforms as T

def get_transform(train):
    transforms = []
    # converts the image, a PIL image, into a PyTorch Tensor
    transforms.append(T.ToTensor())
    if train:
        # during training, randomly flip the training images
        # and ground-truth for data augmentation
        transforms.append(T.RandomHorizontalFlip(0.5))
    return T.Compose(transforms)
```

Putting everything together

We now have the dataset class, the models and the data transforms. Let's instantiate them

```
# use our dataset and defined transformations
dataset = PennFudanDataset('PennFudanPed', get_transform(train=True))
dataset_test = PennFudanDataset('PennFudanPed', get_transform(train=False))

# split the dataset in train and test set
torch.manual_seed(1)
indices = torch.randperm(len(dataset)).tolist()
dataset = torch.utils.data.Subset(dataset, indices[:-50])
dataset_test = torch.utils.data.Subset(dataset_test, indices[-50:])

# define training and validation data loaders
data_loader = torch.utils.data.DataLoader(
    dataset, batch_size=2, shuffle=True, num_workers=4,
```

```
collate_fn=utils.collate_fn)

data_loader_test = torch.utils.data.DataLoader(
    dataset_test, batch_size=1, shuffle=False, num_workers=4,
    collate_fn=utils.collate_fn)
```

Now let's instantiate the model (Option 1: ResNet) and the optimizer

```
device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')

# our dataset has two classes only - background and person
num_classes = 2

# get the model using our helper function
model_1 = get_instance_segmentation_model_1(num_classes)

# move model to the right device
model_1.to(device)

# construct an optimizer
params = [p for p in model_1.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(params, lr=0.005,
                             momentum=0.9, weight_decay=0.0005)

# and a learning rate scheduler which decreases the learning rate by
# 10x every 3 epochs
lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
                                                step_size=3,
                                                gamma=0.1)
```

And now let's train the model for 10 epochs, evaluating at the end of 10th epoch.

```
# let's train it for 10 epochs
from torch.optim.lr_scheduler import StepLR
num_epochs = 10

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    train_one_epoch(model_1, optimizer, data_loader, device, epoch, print_freq=10)
    # update the learning rate
    lr_scheduler.step()
```



```

Epoch: [0] [ 0/60] eta: 0:07:17 lr: 0.000090 loss: 2.7899 (2.7899) loss_classifier: 0.7401 (0.7401)
loss_box_reg: 0.3405 (0.3405) loss_mask: 1.6637 (1.6637) loss_objectness: 0.0430 (0.0430)
loss_rpn_box_reg: 0.0025 (0.0025) time: 7.2903 data: 0.3273 max mem: 2162
Epoch: [0] [10/60] eta: 0:00:59 lr: 0.000936 loss: 1.3949 (1.7270) loss_classifier: 0.5158 (0.4826)
loss_box_reg: 0.2960 (0.2981) loss_mask: 0.7157 (0.9198) loss_objectness: 0.0169 (0.0217)
loss_rpn_box_reg: 0.0045 (0.0048) time: 1.1989 data: 0.0434 max mem: 3320
Epoch: [0] [20/60] eta: 0:00:35 lr: 0.001783 loss: 1.0076 (1.2311) loss_classifier: 0.2244 (0.3358)
loss_box_reg: 0.2910 (0.2865) loss_mask: 0.3239 (0.5872) loss_objectness: 0.0109 (0.0171)
loss_rpn_box_reg: 0.0042 (0.0045) time: 0.5689 data: 0.0180 max mem: 3320
Epoch: [0] [30/60] eta: 0:00:23 lr: 0.002629 loss: 0.5587 (1.0162) loss_classifier: 0.0983 (0.2554)
loss_box_reg: 0.2665 (0.2870) loss_mask: 0.1833 (0.4537) loss_objectness: 0.0102 (0.0151)
loss_rpn_box_reg: 0.0045 (0.0050) time: 0.5637 data: 0.0178 max mem: 3320
Epoch: [0] [40/60] eta: 0:00:14 lr: 0.003476 loss: 0.4513 (0.8823) loss_classifier: 0.0606 (0.2065)
loss_box_reg: 0.2153 (0.2675) loss_mask: 0.1689 (0.3901) loss_objectness: 0.0083 (0.0128)
loss_rpn_box_reg: 0.0056 (0.0053) time: 0.5675 data: 0.0150 max mem: 3320
Epoch: [0] [50/60] eta: 0:00:06 lr: 0.004323 loss: 0.3895 (0.7843) loss_classifier: 0.0418 (0.1747)
loss_box_reg: 0.1625 (0.2453) loss_mask: 0.1720 (0.3475) loss_objectness: 0.0035 (0.0109)
loss_rpn_box_reg: 0.0049 (0.0058) time: 0.5375 data: 0.0131 max mem: 3320
Epoch: [0] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.3133 (0.7138) loss_classifier: 0.0366 (0.1543)
loss_box_reg: 0.1154 (0.2265) loss_mask: 0.1490 (0.3173) loss_objectness: 0.0013 (0.0097)
loss_rpn_box_reg: 0.0059 (0.0060) time: 0.5332 data: 0.0106 max mem: 3320
Epoch: [0] Total time: 0:00:40 (0.6691 s / it)
Epoch: [1] [ 0/60] eta: 0:01:07 lr: 0.005000 loss: 0.3843 (0.3843) loss_classifier: 0.0526 (0.0526)
loss_box_reg: 0.1679 (0.1679) loss_mask: 0.1554 (0.1554) loss_objectness: 0.0003 (0.0003)
loss_rpn_box_reg: 0.0081 (0.0081) time: 1.1258 data: 0.5763 max mem: 3320
Epoch: [1] [10/60] eta: 0:00:29 lr: 0.005000 loss: 0.2921 (0.2999) loss_classifier: 0.0359 (0.0398)
loss_box_reg: 0.1132 (0.0989) loss_mask: 0.1438 (0.1530) loss_objectness: 0.0010 (0.0020)
loss_rpn_box_reg: 0.0058 (0.0063) time: 0.5822 data: 0.0619 max mem: 3320
Epoch: [1] [20/60] eta: 0:00:23 lr: 0.005000 loss: 0.2389 (0.2900) loss_classifier: 0.0358 (0.0390)
loss_box_reg: 0.0766 (0.0950) loss_mask: 0.1369 (0.1479) loss_objectness: 0.0012 (0.0025)
loss_rpn_box_reg: 0.0046 (0.0056) time: 0.5504 data: 0.0108 max mem: 3320

```

```

Epoch: [1] [30/60] eta: 0:00:17 lr: 0.005000 loss: 0.3070 (0.3076) loss_classifier: 0.0380 (0.0421)
loss_box_reg: 0.0953 (0.1044) loss_mask: 0.1385 (0.1524) loss_objectness: 0.0013 (0.0026)
loss_rpn_box_reg: 0.0046 (0.0061) time: 0.5820 data: 0.0124 max mem: 3320
Epoch: [1] [40/60] eta: 0:00:11 lr: 0.005000 loss: 0.3070 (0.3003) loss_classifier: 0.0380 (0.0408)
loss_box_reg: 0.0794 (0.1003) loss_mask: 0.1420 (0.1510) loss_objectness: 0.0006 (0.0024)
loss_rpn_box_reg: 0.0057 (0.0058) time: 0.5662 data: 0.0120 max mem: 3320
Epoch: [1] [50/60] eta: 0:00:05 lr: 0.005000 loss: 0.2550 (0.2947) loss_classifier: 0.0331 (0.0395)
loss_box_reg: 0.0709 (0.0949) loss_mask: 0.1420 (0.1527) loss_objectness: 0.0006 (0.0023)
loss_rpn_box_reg: 0.0037 (0.0053) time: 0.5447 data: 0.0123 max mem: 3320
Epoch: [1] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.3055 (0.2972) loss_classifier: 0.0358 (0.0396)
loss_box_reg: 0.0753 (0.0948) loss_mask: 0.1549 (0.1547) loss_objectness: 0.0011 (0.0025)
loss_rpn_box_reg: 0.0047 (0.0056) time: 0.5559 data: 0.0119 max mem: 3320
Epoch: [1] Total time: 0:00:34 (0.5708 s / it)
Epoch: [2] [ 0/60] eta: 0:00:51 lr: 0.005000 loss: 0.3595 (0.3595) loss_classifier: 0.0293 (0.0293)
loss_box_reg: 0.0862 (0.0862) loss_mask: 0.2358 (0.2358) loss_objectness: 0.0010 (0.0010)
loss_rpn_box_reg: 0.0072 (0.0072) time: 0.8561 data: 0.3386 max mem: 3320
Epoch: [2] [10/60] eta: 0:00:30 lr: 0.005000 loss: 0.2915 (0.2671) loss_classifier: 0.0353 (0.0352)
loss_box_reg: 0.0756 (0.0725) loss_mask: 0.1503 (0.1528) loss_objectness: 0.0010 (0.0017)
loss_rpn_box_reg: 0.0046 (0.0048) time: 0.6186 data: 0.0416 max mem: 3320
Epoch: [2] [20/60] eta: 0:00:24 lr: 0.005000 loss: 0.2858 (0.2811) loss_classifier: 0.0414 (0.0433)
loss_box_reg: 0.0756 (0.0840) loss_mask: 0.1367 (0.1469) loss_objectness: 0.0011 (0.0016)
loss_rpn_box_reg: 0.0048 (0.0053) time: 0.5967 data: 0.0115 max mem: 3320
Epoch: [2] [30/60] eta: 0:00:18 lr: 0.005000 loss: 0.2372 (0.2560) loss_classifier: 0.0373 (0.0380)
loss_box_reg: 0.0632 (0.0735) loss_mask: 0.1367 (0.1385) loss_objectness: 0.0005 (0.0013)
loss_rpn_box_reg: 0.0033 (0.0046) time: 0.5927 data: 0.0133 max mem: 3320
Epoch: [2] [40/60] eta: 0:00:11 lr: 0.005000 loss: 0.1995 (0.2480) loss_classifier: 0.0249 (0.0365)
loss_box_reg: 0.0459 (0.0708) loss_mask: 0.1199 (0.1352) loss_objectness: 0.0003 (0.0012)
loss_rpn_box_reg: 0.0022 (0.0042) time: 0.5795 data: 0.0126 max mem: 3320
Epoch: [2] [50/60] eta: 0:00:05 lr: 0.005000 loss: 0.2178 (0.2453) loss_classifier: 0.0324 (0.0357)
loss_box_reg: 0.0540 (0.0697) loss_mask: 0.1232 (0.1345) loss_objectness: 0.0006 (0.0013)
loss_rpn_box_reg: 0.0026 (0.0041) time: 0.5692 data: 0.0115 max mem: 3320

```

```

Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2120 (0.2412) loss_classifier: 0.0290 (0.0345)
loss_box_reg: 0.0531 (0.0675) loss_mask: 0.1282 (0.1341) loss_objectness: 0.0004 (0.0013)
loss_rpn_box_reg: 0.0022 (0.0039) time: 0.5510 data: 0.0111 max mem: 3320
Epoch: [2] Total time: 0:00:34 (0.5825 s / it)
Epoch: [3] [ 0/60] eta: 0:00:54 lr: 0.000500 loss: 0.2435 (0.2435) loss_classifier: 0.0499 (0.0499)
loss_box_reg: 0.0648 (0.0648) loss_mask: 0.1255 (0.1255) loss_objectness: 0.0023 (0.0023)
loss_rpn_box_reg: 0.0010 (0.0010) time: 0.9023 data: 0.3088 max mem: 3320
Epoch: [3] [10/60] eta: 0:00:31 lr: 0.000500 loss: 0.2206 (0.2460) loss_classifier: 0.0300 (0.0339)
loss_box_reg: 0.0597 (0.0639) loss_mask: 0.1291 (0.1426) loss_objectness: 0.0007 (0.0012)
loss_rpn_box_reg: 0.0043 (0.0045) time: 0.6361 data: 0.0380 max mem: 3320
Epoch: [3] [20/60] eta: 0:00:24 lr: 0.000500 loss: 0.2158 (0.2326) loss_classifier: 0.0297 (0.0328)
loss_box_reg: 0.0597 (0.0616) loss_mask: 0.1202 (0.1333) loss_objectness: 0.0007 (0.0011)
loss_rpn_box_reg: 0.0038 (0.0039) time: 0.5939 data: 0.0110 max mem: 3320
Epoch: [3] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1909 (0.2232) loss_classifier: 0.0293 (0.0305)
loss_box_reg: 0.0534 (0.0570) loss_mask: 0.1128 (0.1306) loss_objectness: 0.0005 (0.0013)
loss_rpn_box_reg: 0.0026 (0.0037) time: 0.5725 data: 0.0108 max mem: 3320
Epoch: [3] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1956 (0.2158) loss_classifier: 0.0233 (0.0294)
loss_box_reg: 0.0470 (0.0546) loss_mask: 0.1174 (0.1271) loss_objectness: 0.0004 (0.0011)
loss_rpn_box_reg: 0.0029 (0.0036) time: 0.5741 data: 0.0106 max mem: 3320
Epoch: [3] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1823 (0.2084) loss_classifier: 0.0233 (0.0292)
loss_box_reg: 0.0374 (0.0512) loss_mask: 0.1132 (0.1235) loss_objectness: 0.0003 (0.0012)
loss_rpn_box_reg: 0.0023 (0.0034) time: 0.5779 data: 0.0101 max mem: 3320
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1814 (0.2100) loss_classifier: 0.0235 (0.0300)
loss_box_reg: 0.0334 (0.0509) loss_mask: 0.1119 (0.1246) loss_objectness: 0.0003 (0.0011)
loss_rpn_box_reg: 0.0020 (0.0034) time: 0.5755 data: 0.0103 max mem: 3320
Epoch: [3] Total time: 0:00:35 (0.5885 s / it)
Epoch: [4] [ 0/60] eta: 0:00:50 lr: 0.000500 loss: 0.1429 (0.1429) loss_classifier: 0.0145 (0.0145)
loss_box_reg: 0.0155 (0.0155) loss_mask: 0.1108 (0.1108) loss_objectness: 0.0003 (0.0003)
loss_rpn_box_reg: 0.0018 (0.0018) time: 0.8337 data: 0.3439 max mem: 3320
Epoch: [4] [10/60] eta: 0:00:29 lr: 0.000500 loss: 0.1837 (0.1830) loss_classifier: 0.0249 (0.0255)
loss_box_reg: 0.0379 (0.0411) loss_mask: 0.1135 (0.1128) loss_objectness: 0.0005 (0.0013)

```

```

loss_rpn_box_reg: 0.0016 (0.0022) time: 0.5834 data: 0.0396 max mem: 3320
Epoch: [4] [20/60] eta: 0:00:23 lr: 0.000500 loss: 0.1867 (0.2062) loss_classifier: 0.0298 (0.0327)
loss_box_reg: 0.0379 (0.0467) loss_mask: 0.1149 (0.1224) loss_objectness: 0.0005 (0.0014)
loss_rpn_box_reg: 0.0016 (0.0031) time: 0.5725 data: 0.0106 max mem: 3320
Epoch: [4] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1867 (0.2043) loss_classifier: 0.0358 (0.0322)
loss_box_reg: 0.0426 (0.0460) loss_mask: 0.1203 (0.1215) loss_objectness: 0.0004 (0.0012)
loss_rpn_box_reg: 0.0028 (0.0034) time: 0.5962 data: 0.0112 max mem: 3320
Epoch: [4] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1772 (0.1971) loss_classifier: 0.0227 (0.0300)
loss_box_reg: 0.0349 (0.0438) loss_mask: 0.1148 (0.1191) loss_objectness: 0.0004 (0.0011)
loss_rpn_box_reg: 0.0020 (0.0032) time: 0.5807 data: 0.0121 max mem: 3320
Epoch: [4] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1575 (0.1912) loss_classifier: 0.0187 (0.0286)
loss_box_reg: 0.0285 (0.0419) loss_mask: 0.1104 (0.1167) loss_objectness: 0.0002 (0.0010)
loss_rpn_box_reg: 0.0019 (0.0030) time: 0.5492 data: 0.0130 max mem: 3320
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1653 (0.1873) loss_classifier: 0.0196 (0.0279)
loss_box_reg: 0.0296 (0.0404) loss_mask: 0.1097 (0.1152) loss_objectness: 0.0002 (0.0010)
loss_rpn_box_reg: 0.0020 (0.0029) time: 0.5691 data: 0.0112 max mem: 3320
Epoch: [4] Total time: 0:00:34 (0.5804 s / it)
Epoch: [5] [ 0/60] eta: 0:00:57 lr: 0.000500 loss: 0.1310 (0.1310) loss_classifier: 0.0131 (0.0131)
loss_box_reg: 0.0213 (0.0213) loss_mask: 0.0950 (0.0950) loss_objectness: 0.0001 (0.0001)
loss_rpn_box_reg: 0.0016 (0.0016) time: 0.9573 data: 0.3209 max mem: 3320
Epoch: [5] [10/60] eta: 0:00:28 lr: 0.000500 loss: 0.1633 (0.1947) loss_classifier: 0.0207 (0.0277)
loss_box_reg: 0.0325 (0.0452) loss_mask: 0.1095 (0.1180) loss_objectness: 0.0005 (0.0016)
loss_rpn_box_reg: 0.0016 (0.0021) time: 0.5658 data: 0.0379 max mem: 3320
Epoch: [5] [20/60] eta: 0:00:23 lr: 0.000500 loss: 0.1681 (0.1876) loss_classifier: 0.0210 (0.0264)
loss_box_reg: 0.0325 (0.0424) loss_mask: 0.1095 (0.1154) loss_objectness: 0.0005 (0.0011)
loss_rpn_box_reg: 0.0018 (0.0022) time: 0.5605 data: 0.0108 max mem: 3502
Epoch: [5] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1723 (0.1936) loss_classifier: 0.0235 (0.0273)
loss_box_reg: 0.0349 (0.0441) loss_mask: 0.1086 (0.1184) loss_objectness: 0.0004 (0.0010)
loss_rpn_box_reg: 0.0025 (0.0029) time: 0.5931 data: 0.0115 max mem: 3502
Epoch: [5] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1674 (0.1893) loss_classifier: 0.0235 (0.0265)
loss_box_reg: 0.0320 (0.0421) loss_mask: 0.1080 (0.1171) loss_objectness: 0.0004 (0.0008)

```

```

loss_rpn_box_reg: 0.0026 (0.0028) time: 0.5918 data: 0.0121 max mem: 3502
Epoch: [5] [50/60] eta: 0:00:05 lr: 0.0000500 loss: 0.1656 (0.1929) loss_classifier: 0.0194 (0.0274)
loss_box_reg: 0.0254 (0.0436) loss_mask: 0.1074 (0.1181) loss_objectness: 0.0002 (0.0011)
loss_rpn_box_reg: 0.0019 (0.0028) time: 0.5911 data: 0.0120 max mem: 3502
Epoch: [5] [59/60] eta: 0:00:00 lr: 0.0000500 loss: 0.1656 (0.1924) loss_classifier: 0.0194 (0.0274)
loss_box_reg: 0.0321 (0.0433) loss_mask: 0.1074 (0.1178) loss_objectness: 0.0003 (0.0010)
loss_rpn_box_reg: 0.0023 (0.0029) time: 0.5882 data: 0.0108 max mem: 3502
Epoch: [5] Total time: 0:00:35 (0.5912 s / it)
Epoch: [6] [ 0/60] eta: 0:01:13 lr: 0.0000500 loss: 0.1655 (0.1655) loss_classifier: 0.0326 (0.0326)
loss_box_reg: 0.0281 (0.0281) loss_mask: 0.0993 (0.0993) loss_objectness: 0.0043 (0.0043)
loss_rpn_box_reg: 0.0012 (0.0012) time: 1.2214 data: 0.6604 max mem: 3502
Epoch: [6] [10/60] eta: 0:00:32 lr: 0.0000500 loss: 0.2183 (0.2191) loss_classifier: 0.0304 (0.0331)
loss_box_reg: 0.0422 (0.0537) loss_mask: 0.1325 (0.1283) loss_objectness: 0.0008 (0.0013)
loss_rpn_box_reg: 0.0031 (0.0027) time: 0.6473 data: 0.0730 max mem: 3502
Epoch: [6] [20/60] eta: 0:00:24 lr: 0.0000500 loss: 0.1709 (0.1974) loss_classifier: 0.0263 (0.0298)
loss_box_reg: 0.0350 (0.0455) loss_mask: 0.1086 (0.1180) loss_objectness: 0.0005 (0.0015)
loss_rpn_box_reg: 0.0026 (0.0026) time: 0.5931 data: 0.0127 max mem: 3502
Epoch: [6] [30/60] eta: 0:00:18 lr: 0.0000500 loss: 0.1611 (0.1851) loss_classifier: 0.0207 (0.0269)
loss_box_reg: 0.0265 (0.0403) loss_mask: 0.1022 (0.1145) loss_objectness: 0.0002 (0.0011)
loss_rpn_box_reg: 0.0018 (0.0023) time: 0.5922 data: 0.0111 max mem: 3502
Epoch: [6] [40/60] eta: 0:00:12 lr: 0.0000500 loss: 0.1757 (0.1900) loss_classifier: 0.0236 (0.0274)
loss_box_reg: 0.0359 (0.0424) loss_mask: 0.1055 (0.1166) loss_objectness: 0.0004 (0.0011)
loss_rpn_box_reg: 0.0019 (0.0025) time: 0.6014 data: 0.0122 max mem: 3502
Epoch: [6] [50/60] eta: 0:00:06 lr: 0.0000500 loss: 0.1819 (0.1885) loss_classifier: 0.0230 (0.0268)
loss_box_reg: 0.0383 (0.0413) loss_mask: 0.1144 (0.1165) loss_objectness: 0.0004 (0.0013)
loss_rpn_box_reg: 0.0026 (0.0026) time: 0.5908 data: 0.0145 max mem: 3502
Epoch: [6] [59/60] eta: 0:00:00 lr: 0.0000500 loss: 0.1558 (0.1880) loss_classifier: 0.0156 (0.0268)
loss_box_reg: 0.0255 (0.0408) loss_mask: 0.1106 (0.1168) loss_objectness: 0.0004 (0.0012)
loss_rpn_box_reg: 0.0021 (0.0025) time: 0.5713 data: 0.0134 max mem: 3502
Epoch: [6] Total time: 0:00:35 (0.5997 s / it)
Epoch: [7] [ 0/60] eta: 0:00:57 lr: 0.0000500 loss: 0.1362 (0.1362) loss_classifier: 0.0213 (0.0213)

```

```

loss_box_reg: 0.0217 (0.0217) loss_mask: 0.0909 (0.0909) loss_objectness: 0.0009 (0.0009)
loss_rpn_box_reg: 0.0014 (0.0014) time: 0.9515 data: 0.3533 max mem: 3502
Epoch: [7] [10/60] eta: 0:00:29 lr: 0.000050 loss: 0.1661 (0.1890) loss_classifier: 0.0232 (0.0290)
loss_box_reg: 0.0328 (0.0385) loss_mask: 0.1108 (0.1188) loss_objectness: 0.0003 (0.0004)
loss_rpn_box_reg: 0.0021 (0.0023) time: 0.5984 data: 0.0443 max mem: 3502
Epoch: [7] [20/60] eta: 0:00:23 lr: 0.000050 loss: 0.1684 (0.1928) loss_classifier: 0.0232 (0.0273)
loss_box_reg: 0.0328 (0.0403) loss_mask: 0.1180 (0.1222) loss_objectness: 0.0002 (0.0003)
loss_rpn_box_reg: 0.0021 (0.0027) time: 0.5674 data: 0.0118 max mem: 3502
Epoch: [7] [30/60] eta: 0:00:17 lr: 0.000050 loss: 0.1791 (0.1899) loss_classifier: 0.0228 (0.0278)
loss_box_reg: 0.0301 (0.0408) loss_mask: 0.1161 (0.1181) loss_objectness: 0.0002 (0.0006)
loss_rpn_box_reg: 0.0019 (0.0026) time: 0.5950 data: 0.0120 max mem: 3502
Epoch: [7] [40/60] eta: 0:00:11 lr: 0.000050 loss: 0.1677 (0.1861) loss_classifier: 0.0231 (0.0268)
loss_box_reg: 0.0338 (0.0395) loss_mask: 0.1100 (0.1165) loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0026 (0.0026) time: 0.6054 data: 0.0122 max mem: 3502
Epoch: [7] [50/60] eta: 0:00:05 lr: 0.000050 loss: 0.1652 (0.1857) loss_classifier: 0.0235 (0.0266)
loss_box_reg: 0.0341 (0.0396) loss_mask: 0.1062 (0.1161) loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0026 (0.0027) time: 0.5934 data: 0.0117 max mem: 3502
Epoch: [7] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1715 (0.1867) loss_classifier: 0.0270 (0.0267)
loss_box_reg: 0.0329 (0.0398) loss_mask: 0.1116 (0.1167) loss_objectness: 0.0002 (0.0007)
loss_rpn_box_reg: 0.0019 (0.0026) time: 0.5712 data: 0.0110 max mem: 3502
Epoch: [7] Total time: 0:00:35 (0.5880 s / it)
Epoch: [8] [ 0/60] eta: 0:00:59 lr: 0.000050 loss: 0.2865 (0.2865) loss_classifier: 0.0580 (0.0580)
loss_box_reg: 0.0720 (0.0720) loss_mask: 0.1505 (0.1505) loss_objectness: 0.0012 (0.0012)
loss_rpn_box_reg: 0.0049 (0.0049) time: 0.9995 data: 0.3814 max mem: 3502
Epoch: [8] [10/60] eta: 0:00:30 lr: 0.000050 loss: 0.1997 (0.2095) loss_classifier: 0.0277 (0.0314)
loss_box_reg: 0.0463 (0.0488) loss_mask: 0.1209 (0.1248) loss_objectness: 0.0004 (0.0007)
loss_rpn_box_reg: 0.0034 (0.0038) time: 0.6164 data: 0.0449 max mem: 3502
Epoch: [8] [20/60] eta: 0:00:24 lr: 0.000050 loss: 0.1850 (0.2041) loss_classifier: 0.0279 (0.0306)
loss_box_reg: 0.0408 (0.0452) loss_mask: 0.1157 (0.1247) loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0022 (0.0029) time: 0.5852 data: 0.0110 max mem: 3502
Epoch: [8] [30/60] eta: 0:00:18 lr: 0.000050 loss: 0.1732 (0.2025) loss_classifier: 0.0288 (0.0301)

```

```

loss_box_reg: 0.0408 (0.0461) loss_mask: 0.1126 (0.1224) loss_objectness: 0.0003 (0.0009)
loss_rpn_box_reg: 0.0020 (0.0030) time: 0.5965 data: 0.0125 max mem: 3502
Epoch: [8] [40/60] eta: 0:00:11 lr: 0.000050 loss: 0.1583 (0.1962) loss_classifier: 0.0219 (0.0287)
loss_box_reg: 0.0342 (0.0437) loss_mask: 0.1055 (0.1202) loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0024 (0.0029) time: 0.5761 data: 0.0120 max mem: 3502
Epoch: [8] [50/60] eta: 0:00:05 lr: 0.000050 loss: 0.1523 (0.1886) loss_classifier: 0.0207 (0.0272)
loss_box_reg: 0.0241 (0.0406) loss_mask: 0.1015 (0.1174) loss_objectness: 0.0002 (0.0007)
loss_rpn_box_reg: 0.0012 (0.0027) time: 0.5625 data: 0.0106 max mem: 3502
Epoch: [8] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1590 (0.1897) loss_classifier: 0.0245 (0.0272)
loss_box_reg: 0.0258 (0.0407) loss_mask: 0.1005 (0.1184) loss_objectness: 0.0002 (0.0006)
loss_rpn_box_reg: 0.0017 (0.0028) time: 0.5649 data: 0.0110 max mem: 3502
Epoch: [8] Total time: 0:00:35 (0.5856 s / it)
Epoch: [9] [ 0/60] eta: 0:00:51 lr: 0.000005 loss: 0.1334 (0.1334) loss_classifier: 0.0115 (0.0115)
loss_box_reg: 0.0185 (0.0185) loss_mask: 0.1027 (0.1027) loss_objectness: 0.0003 (0.0003)
loss_rpn_box_reg: 0.0003 (0.0003) time: 0.8519 data: 0.2903 max mem: 3502
Epoch: [9] [10/60] eta: 0:00:30 lr: 0.000005 loss: 0.1661 (0.1952) loss_classifier: 0.0255 (0.0257)
loss_box_reg: 0.0370 (0.0407) loss_mask: 0.1086 (0.1237) loss_objectness: 0.0003 (0.0011)
loss_rpn_box_reg: 0.0044 (0.0039) time: 0.6142 data: 0.0361 max mem: 3502
Epoch: [9] [20/60] eta: 0:00:24 lr: 0.000005 loss: 0.1924 (0.2076) loss_classifier: 0.0308 (0.0292)
loss_box_reg: 0.0454 (0.0485) loss_mask: 0.1188 (0.1246) loss_objectness: 0.0006 (0.0015)
loss_rpn_box_reg: 0.0029 (0.0038) time: 0.6057 data: 0.0130 max mem: 3502
Epoch: [9] [30/60] eta: 0:00:18 lr: 0.000005 loss: 0.1712 (0.1939) loss_classifier: 0.0256 (0.0265)
loss_box_reg: 0.0339 (0.0422) loss_mask: 0.1153 (0.1209) loss_objectness: 0.0004 (0.0012)
loss_rpn_box_reg: 0.0022 (0.0031) time: 0.5977 data: 0.0129 max mem: 3601
Epoch: [9] [40/60] eta: 0:00:11 lr: 0.000005 loss: 0.1612 (0.1870) loss_classifier: 0.0205 (0.0261)
loss_box_reg: 0.0229 (0.0386) loss_mask: 0.1069 (0.1183) loss_objectness: 0.0002 (0.0011)
loss_rpn_box_reg: 0.0015 (0.0029) time: 0.5673 data: 0.0123 max mem: 3601
Epoch: [9] [50/60] eta: 0:00:05 lr: 0.000005 loss: 0.1615 (0.1865) loss_classifier: 0.0239 (0.0272)
loss_box_reg: 0.0288 (0.0387) loss_mask: 0.1023 (0.1167) loss_objectness: 0.0004 (0.0012)
loss_rpn_box_reg: 0.0014 (0.0027) time: 0.5674 data: 0.0122 max mem: 3601
Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.1733 (0.1867) loss_classifier: 0.0238 (0.0275)

```

```
loss_box_reg: 0.0291 (0.0389) loss_mask: 0.1071 (0.1163) loss_objectness: 0.0004 (0.0012)
loss_rpn_box_reg: 0.0014 (0.0028) time: 0.5804 data: 0.0106 max mem: 3601
Epoch: [9] Total time: 0:00:35 (0.5911 s / it)
```


Evaluate the model (Option 1: ResNet)

```
evaluate(model_1, data_loader_test, device=device)
```

```
creating index...
```

```
index created!
```

```
Test: [ 0/50] eta: 0:00:19 model_time: 0.1760 (0.1760) evaluator_time: 0.0051 (0.0051)
      time: 0.3944 data: 0.2121 max mem: 3601
```

```
Test: [49/50] eta: 0:00:00 model_time: 0.1104 (0.1096) evaluator_time: 0.0046 (0.0062)
      time: 0.1239 data: 0.0066 max mem: 3601
```

```
Test: Total time: 0:00:06 (0.1300 s / it)
```

```
Averaged stats: model_time: 0.1104 (0.1096) evaluator_time: 0.0046 (0.0062)
```

```
Accumulating evaluation results...
```

```
DONE (t=0.02s).
```

```
Accumulating evaluation results...
```

```
DONE (t=0.02s).
```

```
IoU metric: bbox
```

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.834
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.990
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.960
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.508
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.846
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.380
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.876
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.876
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.775
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.883
```

```
IoU metric: segm
```

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.759
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.990
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.918
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.452
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.770
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.345
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.802
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.802
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.725
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.807
```

Now let's instantiate the model (Option 2: MobileNet) and the optimizer.

```
# get the model using our helper function
model_2 = get_instance_segmentation_model_2(num_classes)

# move model to the right device
model_2.to(device)

# construct an optimizer
params = [p for p in model_2.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(params, lr=0.005,
                             momentum=0.9, weight_decay=0.0005)

# and a learning rate scheduler which decreases the learning rate by
# 10x every 3 epochs
lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
                                                step_size=3,
                                                gamma=0.1)
```

And now let's train the model for 10 epochs, evaluating at the end of 10th epoch.

```
# let's train it for 10 epochs
from torch.optim.lr_scheduler import StepLR
num_epochs = 10

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    train_one_epoch(model_2, optimizer, data_loader, device, epoch, print_freq=10)
    # update the learning rate
    lr_scheduler.step()
```

```

Epoch: [0] [ 0/60] eta: 0:01:05 lr: 0.000090 loss: 6.5203 (6.5203) loss_classifier: 0.6889 (0.6889)
loss_box_reg: 0.0088 (0.0088) loss_mask: 5.0657 (5.0657) loss_objectness: 0.6994 (0.6994)
loss_rpn_box_reg: 0.0575 (0.0575) time: 1.0968 data: 0.2465 max mem: 3693
Epoch: [0] [10/60] eta: 0:00:21 lr: 0.000936 loss: 4.7205 (4.6260) loss_classifier: 0.3987 (0.3992)
loss_box_reg: 0.0289 (0.0326) loss_mask: 3.6863 (3.4848) loss_objectness: 0.6767 (0.6646)
loss_rpn_box_reg: 0.0352 (0.0449) time: 0.4322 data: 0.0316 max mem: 5262
Epoch: [0] [20/60] eta: 0:00:16 lr: 0.001783 loss: 2.6217 (3.3756) loss_classifier: 0.2578 (0.3371)
loss_box_reg: 0.0701 (0.0947) loss_mask: 1.6175 (2.3442) loss_objectness: 0.5531 (0.5599)
loss_rpn_box_reg: 0.0245 (0.0397) time: 0.3894 data: 0.0102 max mem: 5724
Epoch: [0] [30/60] eta: 0:00:12 lr: 0.002629 loss: 1.5354 (2.7402) loss_classifier: 0.2467 (0.3286)
loss_box_reg: 0.1486 (0.1191) loss_mask: 0.6964 (1.7874) loss_objectness: 0.3265 (0.4687)
loss_rpn_box_reg: 0.0254 (0.0363) time: 0.4320 data: 0.0132 max mem: 6265
Epoch: [0] [40/60] eta: 0:00:08 lr: 0.003476 loss: 1.2901 (2.3946) loss_classifier: 0.2437 (0.3174)
loss_box_reg: 0.1729 (0.1362) loss_mask: 0.6119 (1.5038) loss_objectness: 0.2244 (0.4024)
loss_rpn_box_reg: 0.0261 (0.0347) time: 0.4314 data: 0.0137 max mem: 6265
Epoch: [0] [50/60] eta: 0:00:04 lr: 0.004323 loss: 1.0495 (2.1153) loss_classifier: 0.1992 (0.2912)
loss_box_reg: 0.1809 (0.1415) loss_mask: 0.5276 (1.3044) loss_objectness: 0.1576 (0.3460)
loss_rpn_box_reg: 0.0220 (0.0323) time: 0.4102 data: 0.0104 max mem: 6265
Epoch: [0] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.9402 (1.9310) loss_classifier: 0.1703 (0.2699)
loss_box_reg: 0.1507 (0.1416) loss_mask: 0.4977 (1.1843) loss_objectness: 0.0928 (0.3037)
loss_rpn_box_reg: 0.0220 (0.0314) time: 0.3988 data: 0.0101 max mem: 6265
Epoch: [0] Total time: 0:00:25 (0.4208 s / it)
Epoch: [1] [ 0/60] eta: 0:00:46 lr: 0.005000 loss: 0.7883 (0.7883) loss_classifier: 0.1263 (0.1263)
loss_box_reg: 0.1246 (0.1246) loss_mask: 0.3986 (0.3986) loss_objectness: 0.0752 (0.0752)
loss_rpn_box_reg: 0.0636 (0.0636) time: 0.7824 data: 0.3804 max mem: 6265
Epoch: [1] [10/60] eta: 0:00:21 lr: 0.005000 loss: 0.7559 (0.7262) loss_classifier: 0.1164 (0.1175)
loss_box_reg: 0.1168 (0.1245) loss_mask: 0.3986 (0.3876) loss_objectness: 0.0550 (0.0532)
loss_rpn_box_reg: 0.0359 (0.0433) time: 0.4286 data: 0.0415 max mem: 6265
Epoch: [1] [20/60] eta: 0:00:16 lr: 0.005000 loss: 0.7559 (0.7564) loss_classifier: 0.1161 (0.1238)
loss_box_reg: 0.1168 (0.1292) loss_mask: 0.3963 (0.3975) loss_objectness: 0.0549 (0.0569)
loss_rpn_box_reg: 0.0373 (0.0490) time: 0.4060 data: 0.0089 max mem: 6265

```

```

Epoch: [1] [30/60] eta: 0:00:12 lr: 0.005000 loss: 0.6814 (0.7417) loss_classifier: 0.1022 (0.1179)
loss_box_reg: 0.1128 (0.1267) loss_mask: 0.3963 (0.3971) loss_objectness: 0.0494 (0.0532)
loss_rpn_box_reg: 0.0368 (0.0467) time: 0.4115 data: 0.0132 max mem: 6265
Epoch: [1] [40/60] eta: 0:00:08 lr: 0.005000 loss: 0.6708 (0.7500) loss_classifier: 0.0879 (0.1127)
loss_box_reg: 0.1205 (0.1293) loss_mask: 0.4019 (0.4090) loss_objectness: 0.0428 (0.0510)
loss_rpn_box_reg: 0.0344 (0.0480) time: 0.4071 data: 0.0137 max mem: 6265
Epoch: [1] [50/60] eta: 0:00:04 lr: 0.005000 loss: 0.6906 (0.7310) loss_classifier: 0.0879 (0.1058)
loss_box_reg: 0.1198 (0.1218) loss_mask: 0.3924 (0.4079) loss_objectness: 0.0366 (0.0504)
loss_rpn_box_reg: 0.0344 (0.0451) time: 0.4002 data: 0.0120 max mem: 6265
Epoch: [1] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.6447 (0.7225) loss_classifier: 0.0808 (0.1027)
loss_box_reg: 0.1198 (0.1254) loss_mask: 0.3478 (0.4033) loss_objectness: 0.0321 (0.0483)
loss_rpn_box_reg: 0.0320 (0.0427) time: 0.4088 data: 0.0128 max mem: 6265
Epoch: [1] Total time: 0:00:24 (0.4160 s / it)
Epoch: [2] [ 0/60] eta: 0:00:39 lr: 0.005000 loss: 0.5779 (0.5779) loss_classifier: 0.0611 (0.0611)
loss_box_reg: 0.0954 (0.0954) loss_mask: 0.3804 (0.3804) loss_objectness: 0.0189 (0.0189)
loss_rpn_box_reg: 0.0221 (0.0221) time: 0.6631 data: 0.2427 max mem: 6265
Epoch: [2] [10/60] eta: 0:00:22 lr: 0.005000 loss: 0.6993 (0.6658) loss_classifier: 0.0873 (0.0937)
loss_box_reg: 0.1159 (0.1298) loss_mask: 0.3804 (0.3670) loss_objectness: 0.0348 (0.0413)
loss_rpn_box_reg: 0.0323 (0.0341) time: 0.4453 data: 0.0312 max mem: 6368
Epoch: [2] [20/60] eta: 0:00:17 lr: 0.005000 loss: 0.5693 (0.6020) loss_classifier: 0.0727 (0.0824)
loss_box_reg: 0.1030 (0.1183) loss_mask: 0.3188 (0.3357) loss_objectness: 0.0347 (0.0355)
loss_rpn_box_reg: 0.0259 (0.0300) time: 0.4172 data: 0.0101 max mem: 6368
Epoch: [2] [30/60] eta: 0:00:12 lr: 0.005000 loss: 0.5262 (0.6008) loss_classifier: 0.0684 (0.0802)
loss_box_reg: 0.0967 (0.1210) loss_mask: 0.3032 (0.3357) loss_objectness: 0.0293 (0.0339)
loss_rpn_box_reg: 0.0246 (0.0300) time: 0.4152 data: 0.0121 max mem: 6368
Epoch: [2] [40/60] eta: 0:00:08 lr: 0.005000 loss: 0.5486 (0.6025) loss_classifier: 0.0691 (0.0778)
loss_box_reg: 0.1041 (0.1192) loss_mask: 0.3394 (0.3415) loss_objectness: 0.0303 (0.0347)
loss_rpn_box_reg: 0.0246 (0.0292) time: 0.4140 data: 0.0124 max mem: 6368
Epoch: [2] [50/60] eta: 0:00:04 lr: 0.005000 loss: 0.5385 (0.5909) loss_classifier: 0.0642 (0.0767)
loss_box_reg: 0.0951 (0.1167) loss_mask: 0.3180 (0.3340) loss_objectness: 0.0334 (0.0345)
loss_rpn_box_reg: 0.0221 (0.0290) time: 0.4067 data: 0.0103 max mem: 6368

```

```

Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.4718 (0.5893) loss_classifier: 0.0654 (0.0767)
loss_box_reg: 0.0712 (0.1134) loss_mask: 0.2928 (0.3354) loss_objectness: 0.0289 (0.0339)
loss_rpn_box_reg: 0.0264 (0.0299) time: 0.4024 data: 0.0104 max mem: 6368
Epoch: [2] Total time: 0:00:25 (0.4179 s / it)
Epoch: [3] [ 0/60] eta: 0:00:48 lr: 0.000500 loss: 0.4911 (0.4911) loss_classifier: 0.0379 (0.0379)
loss_box_reg: 0.0909 (0.0909) loss_mask: 0.3092 (0.3092) loss_objectness: 0.0176 (0.0176)
loss_rpn_box_reg: 0.0355 (0.0355) time: 0.8160 data: 0.3878 max mem: 6368
Epoch: [3] [10/60] eta: 0:00:22 lr: 0.000500 loss: 0.4911 (0.5296) loss_classifier: 0.0602 (0.0648)
loss_box_reg: 0.1060 (0.1102) loss_mask: 0.2897 (0.2908) loss_objectness: 0.0255 (0.0283)
loss_rpn_box_reg: 0.0355 (0.0355) time: 0.4529 data: 0.0444 max mem: 6368
Epoch: [3] [20/60] eta: 0:00:17 lr: 0.000500 loss: 0.5026 (0.5378) loss_classifier: 0.0602 (0.0671)
loss_box_reg: 0.0975 (0.1103) loss_mask: 0.2814 (0.2987) loss_objectness: 0.0259 (0.0318)
loss_rpn_box_reg: 0.0302 (0.0298) time: 0.4139 data: 0.0107 max mem: 6368
Epoch: [3] [30/60] eta: 0:00:12 lr: 0.000500 loss: 0.5142 (0.5310) loss_classifier: 0.0598 (0.0677)
loss_box_reg: 0.0889 (0.1087) loss_mask: 0.2754 (0.2980) loss_objectness: 0.0270 (0.0306)
loss_rpn_box_reg: 0.0216 (0.0259) time: 0.4158 data: 0.0136 max mem: 6368
Epoch: [3] [40/60] eta: 0:00:08 lr: 0.000500 loss: 0.5317 (0.5300) loss_classifier: 0.0677 (0.0680)
loss_box_reg: 0.1121 (0.1105) loss_mask: 0.2721 (0.2980) loss_objectness: 0.0220 (0.0292)
loss_rpn_box_reg: 0.0169 (0.0243) time: 0.4232 data: 0.0128 max mem: 6368
Epoch: [3] [50/60] eta: 0:00:04 lr: 0.000500 loss: 0.4661 (0.5222) loss_classifier: 0.0666 (0.0673)
loss_box_reg: 0.1014 (0.1087) loss_mask: 0.2741 (0.2957) loss_objectness: 0.0196 (0.0280)
loss_rpn_box_reg: 0.0158 (0.0225) time: 0.4169 data: 0.0100 max mem: 6368
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.4759 (0.5158) loss_classifier: 0.0616 (0.0667)
loss_box_reg: 0.0944 (0.1087) loss_mask: 0.2616 (0.2901) loss_objectness: 0.0266 (0.0286)
loss_rpn_box_reg: 0.0160 (0.0217) time: 0.4142 data: 0.0104 max mem: 6368
Epoch: [3] Total time: 0:00:25 (0.4259 s / it)
Epoch: [4] [ 0/60] eta: 0:00:43 lr: 0.000500 loss: 0.3010 (0.3010) loss_classifier: 0.0314 (0.0314)
loss_box_reg: 0.0476 (0.0476) loss_mask: 0.1832 (0.1832) loss_objectness: 0.0346 (0.0346)
loss_rpn_box_reg: 0.0042 (0.0042) time: 0.7263 data: 0.2824 max mem: 6368
Epoch: [4] [10/60] eta: 0:00:21 lr: 0.000500 loss: 0.4187 (0.4983) loss_classifier: 0.0492 (0.0556)
loss_box_reg: 0.0775 (0.0988) loss_mask: 0.2722 (0.3019) loss_objectness: 0.0230 (0.0269)

```

```

loss_rpn_box_reg: 0.0118 (0.0151) time: 0.4382 data: 0.0351 max mem: 6368
Epoch: [4] [20/60] eta: 0:00:17 lr: 0.000500 loss: 0.4466 (0.4932) loss_classifier: 0.0567 (0.0602)
loss_box_reg: 0.0989 (0.1112) loss_mask: 0.2742 (0.2837) loss_objectness: 0.0161 (0.0225)
loss_rpn_box_reg: 0.0136 (0.0156) time: 0.4196 data: 0.0104 max mem: 6368
Epoch: [4] [30/60] eta: 0:00:12 lr: 0.000500 loss: 0.4962 (0.4995) loss_classifier: 0.0616 (0.0615)
loss_box_reg: 0.1056 (0.1104) loss_mask: 0.2745 (0.2856) loss_objectness: 0.0184 (0.0264)
loss_rpn_box_reg: 0.0140 (0.0156) time: 0.4253 data: 0.0118 max mem: 6368
Epoch: [4] [40/60] eta: 0:00:08 lr: 0.000500 loss: 0.4710 (0.4949) loss_classifier: 0.0586 (0.0625)
loss_box_reg: 0.0913 (0.1102) loss_mask: 0.2623 (0.2793) loss_objectness: 0.0268 (0.0272)
loss_rpn_box_reg: 0.0145 (0.0157) time: 0.4139 data: 0.0115 max mem: 6368
Epoch: [4] [50/60] eta: 0:00:04 lr: 0.000500 loss: 0.4934 (0.5086) loss_classifier: 0.0725 (0.0661)
loss_box_reg: 0.1135 (0.1151) loss_mask: 0.2645 (0.2837) loss_objectness: 0.0270 (0.0271)
loss_rpn_box_reg: 0.0174 (0.0166) time: 0.4203 data: 0.0100 max mem: 6368
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.4700 (0.4996) loss_classifier: 0.0607 (0.0641)
loss_box_reg: 0.0973 (0.1102) loss_mask: 0.2639 (0.2817) loss_objectness: 0.0240 (0.0270)
loss_rpn_box_reg: 0.0185 (0.0165) time: 0.4144 data: 0.0109 max mem: 6368
Epoch: [4] Total time: 0:00:25 (0.4244 s / it)
Epoch: [5] [ 0/60] eta: 0:00:47 lr: 0.000500 loss: 0.6510 (0.6510) loss_classifier: 0.0944 (0.0944)
loss_box_reg: 0.1744 (0.1744) loss_mask: 0.3214 (0.3214) loss_objectness: 0.0378 (0.0378)
loss_rpn_box_reg: 0.0230 (0.0230) time: 0.7955 data: 0.2694 max mem: 6368
Epoch: [5] [10/60] eta: 0:00:23 lr: 0.000500 loss: 0.6000 (0.6002) loss_classifier: 0.0740 (0.0825)
loss_box_reg: 0.1513 (0.1520) loss_mask: 0.3200 (0.3163) loss_objectness: 0.0251 (0.0270)
loss_rpn_box_reg: 0.0197 (0.0225) time: 0.4691 data: 0.0342 max mem: 6368
Epoch: [5] [20/60] eta: 0:00:17 lr: 0.000500 loss: 0.4647 (0.5521) loss_classifier: 0.0652 (0.0739)
loss_box_reg: 0.1036 (0.1313) loss_mask: 0.2695 (0.3030) loss_objectness: 0.0251 (0.0253)
loss_rpn_box_reg: 0.0157 (0.0185) time: 0.4220 data: 0.0104 max mem: 6368
Epoch: [5] [30/60] eta: 0:00:13 lr: 0.000500 loss: 0.4639 (0.5414) loss_classifier: 0.0635 (0.0712)
loss_box_reg: 0.0942 (0.1242) loss_mask: 0.2873 (0.3011) loss_objectness: 0.0206 (0.0265)
loss_rpn_box_reg: 0.0140 (0.0185) time: 0.4181 data: 0.0131 max mem: 6368
Epoch: [5] [40/60] eta: 0:00:08 lr: 0.000500 loss: 0.4536 (0.5163) loss_classifier: 0.0610 (0.0680)
loss_box_reg: 0.0942 (0.1185) loss_mask: 0.2631 (0.2875) loss_objectness: 0.0166 (0.0250)

```

```

loss_rpn_box_reg: 0.0160 (0.0173) time: 0.4228 data: 0.0128 max mem: 6368
Epoch: [5] [50/60] eta: 0:00:04 lr: 0.000500 loss: 0.4536 (0.5046) loss_classifier: 0.0550 (0.0660)
loss_box_reg: 0.0980 (0.1129) loss_mask: 0.2425 (0.2835) loss_objectness: 0.0166 (0.0249)
loss_rpn_box_reg: 0.0137 (0.0172) time: 0.4131 data: 0.0101 max mem: 6368
Epoch: [5] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.4460 (0.4953) loss_classifier: 0.0523 (0.0643)
loss_box_reg: 0.0980 (0.1105) loss_mask: 0.2582 (0.2804) loss_objectness: 0.0181 (0.0235)
loss_rpn_box_reg: 0.0147 (0.0166) time: 0.4101 data: 0.0110 max mem: 6368
Epoch: [5] Total time: 0:00:25 (0.4258 s / it)
Epoch: [6] [ 0/60] eta: 0:00:52 lr: 0.000050 loss: 0.5759 (0.5759) loss_classifier: 0.0769 (0.0769)
loss_box_reg: 0.0943 (0.0943) loss_mask: 0.3320 (0.3320) loss_objectness: 0.0525 (0.0525)
loss_rpn_box_reg: 0.0202 (0.0202) time: 0.8732 data: 0.4139 max mem: 6368
Epoch: [6] [10/60] eta: 0:00:22 lr: 0.000050 loss: 0.4275 (0.4858) loss_classifier: 0.0617 (0.0650)
loss_box_reg: 0.0877 (0.1023) loss_mask: 0.2712 (0.2770) loss_objectness: 0.0219 (0.0259)
loss_rpn_box_reg: 0.0127 (0.0155) time: 0.4511 data: 0.0446 max mem: 6368
Epoch: [6] [20/60] eta: 0:00:17 lr: 0.000050 loss: 0.4976 (0.5114) loss_classifier: 0.0680 (0.0707)
loss_box_reg: 0.1021 (0.1109) loss_mask: 0.2766 (0.2875) loss_objectness: 0.0209 (0.0255)
loss_rpn_box_reg: 0.0156 (0.0167) time: 0.4148 data: 0.0089 max mem: 6368
Epoch: [6] [30/60] eta: 0:00:13 lr: 0.000050 loss: 0.5398 (0.5245) loss_classifier: 0.0816 (0.0727)
loss_box_reg: 0.1180 (0.1138) loss_mask: 0.2892 (0.2923) loss_objectness: 0.0257 (0.0294)
loss_rpn_box_reg: 0.0161 (0.0163) time: 0.4282 data: 0.0128 max mem: 6368
Epoch: [6] [40/60] eta: 0:00:08 lr: 0.000050 loss: 0.4732 (0.5100) loss_classifier: 0.0675 (0.0696)
loss_box_reg: 0.0876 (0.1101) loss_mask: 0.2790 (0.2863) loss_objectness: 0.0210 (0.0277)
loss_rpn_box_reg: 0.0149 (0.0163) time: 0.4216 data: 0.0131 max mem: 6368
Epoch: [6] [50/60] eta: 0:00:04 lr: 0.000050 loss: 0.4599 (0.5037) loss_classifier: 0.0620 (0.0694)
loss_box_reg: 0.0991 (0.1108) loss_mask: 0.2568 (0.2803) loss_objectness: 0.0210 (0.0267)
loss_rpn_box_reg: 0.0153 (0.0165) time: 0.4163 data: 0.0102 max mem: 6368
Epoch: [6] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.4204 (0.4877) loss_classifier: 0.0586 (0.0670)
loss_box_reg: 0.0933 (0.1073) loss_mask: 0.2284 (0.2729) loss_objectness: 0.0127 (0.0246)
loss_rpn_box_reg: 0.0127 (0.0159) time: 0.4146 data: 0.0105 max mem: 6368
Epoch: [6] Total time: 0:00:25 (0.4272 s / it)
Epoch: [7] [ 0/60] eta: 0:00:47 lr: 0.000050 loss: 0.6768 (0.6768) loss_classifier: 0.0698 (0.0698)

```

```

loss_box_reg: 0.1491 (0.1491) loss_mask: 0.3956 (0.3956) loss_objectness: 0.0404 (0.0404)
loss_rpn_box_reg: 0.0219 (0.0219) time: 0.7872 data: 0.3537 max mem: 6368
Epoch: [7] [10/60] eta: 0:00:22 lr: 0.000050 loss: 0.4558 (0.5038) loss_classifier: 0.0560 (0.0591)
loss_box_reg: 0.1125 (0.1134) loss_mask: 0.2634 (0.2896) loss_objectness: 0.0229 (0.0263)
loss_rpn_box_reg: 0.0177 (0.0153) time: 0.4417 data: 0.0403 max mem: 6368
Epoch: [7] [20/60] eta: 0:00:17 lr: 0.000050 loss: 0.4571 (0.5102) loss_classifier: 0.0609 (0.0649)
loss_box_reg: 0.1125 (0.1206) loss_mask: 0.2608 (0.2846) loss_objectness: 0.0209 (0.0246)
loss_rpn_box_reg: 0.0140 (0.0156) time: 0.4193 data: 0.0104 max mem: 6368
Epoch: [7] [30/60] eta: 0:00:13 lr: 0.000050 loss: 0.4571 (0.4965) loss_classifier: 0.0609 (0.0645)
loss_box_reg: 0.1012 (0.1124) loss_mask: 0.2607 (0.2784) loss_objectness: 0.0178 (0.0255)
loss_rpn_box_reg: 0.0127 (0.0156) time: 0.4288 data: 0.0138 max mem: 6368
Epoch: [7] [40/60] eta: 0:00:08 lr: 0.000050 loss: 0.4190 (0.4798) loss_classifier: 0.0553 (0.0617)
loss_box_reg: 0.0843 (0.1073) loss_mask: 0.2428 (0.2716) loss_objectness: 0.0161 (0.0238)
loss_rpn_box_reg: 0.0123 (0.0153) time: 0.4207 data: 0.0130 max mem: 6368
Epoch: [7] [50/60] eta: 0:00:04 lr: 0.000050 loss: 0.4190 (0.4785) loss_classifier: 0.0551 (0.0631)
loss_box_reg: 0.0856 (0.1066) loss_mask: 0.2387 (0.2702) loss_objectness: 0.0188 (0.0235)
loss_rpn_box_reg: 0.0125 (0.0151) time: 0.4141 data: 0.0101 max mem: 6368
Epoch: [7] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.4685 (0.4824) loss_classifier: 0.0655 (0.0637)
loss_box_reg: 0.0901 (0.1066) loss_mask: 0.2520 (0.2733) loss_objectness: 0.0221 (0.0239)
loss_rpn_box_reg: 0.0122 (0.0148) time: 0.4190 data: 0.0104 max mem: 6368
Epoch: [7] Total time: 0:00:25 (0.4267 s / it)
Epoch: [8] [ 0/60] eta: 0:00:51 lr: 0.000050 loss: 0.6706 (0.6706) loss_classifier: 0.0777 (0.0777)
loss_box_reg: 0.1448 (0.1448) loss_mask: 0.4002 (0.4002) loss_objectness: 0.0295 (0.0295)
loss_rpn_box_reg: 0.0184 (0.0184) time: 0.8535 data: 0.4206 max mem: 6368
Epoch: [8] [10/60] eta: 0:00:23 lr: 0.000050 loss: 0.5050 (0.5125) loss_classifier: 0.0693 (0.0705)
loss_box_reg: 0.1053 (0.1213) loss_mask: 0.2656 (0.2845) loss_objectness: 0.0216 (0.0211)
loss_rpn_box_reg: 0.0146 (0.0151) time: 0.4670 data: 0.0438 max mem: 6368
Epoch: [8] [20/60] eta: 0:00:17 lr: 0.000050 loss: 0.4650 (0.5070) loss_classifier: 0.0617 (0.0686)
loss_box_reg: 0.0953 (0.1107) loss_mask: 0.2535 (0.2800) loss_objectness: 0.0260 (0.0303)
loss_rpn_box_reg: 0.0151 (0.0174) time: 0.4206 data: 0.0081 max mem: 6368
Epoch: [8] [30/60] eta: 0:00:13 lr: 0.000050 loss: 0.4650 (0.5035) loss_classifier: 0.0625 (0.0670)

```



```

loss_box_reg: 0.0953 (0.1113) loss_mask: 0.2566 (0.2809) loss_objectness: 0.0261 (0.0275)
loss_rpn_box_reg: 0.0174 (0.0169) time: 0.4195 data: 0.0116 max mem: 6368
Epoch: [8] [40/60] eta: 0:00:08 lr: 0.0000050 loss: 0.4551 (0.4893) loss_classifier: 0.0607 (0.0647)
loss_box_reg: 0.0977 (0.1070) loss_mask: 0.2736 (0.2753) loss_objectness: 0.0209 (0.0261)
loss_rpn_box_reg: 0.0135 (0.0161) time: 0.4188 data: 0.0114 max mem: 6368
Epoch: [8] [50/60] eta: 0:00:04 lr: 0.0000050 loss: 0.4350 (0.4874) loss_classifier: 0.0607 (0.0639)
loss_box_reg: 0.0896 (0.1066) loss_mask: 0.2736 (0.2770) loss_objectness: 0.0152 (0.0245)
loss_rpn_box_reg: 0.0119 (0.0154) time: 0.4115 data: 0.0101 max mem: 6368
Epoch: [8] [59/60] eta: 0:00:00 lr: 0.0000050 loss: 0.4047 (0.4767) loss_classifier: 0.0543 (0.0628)
loss_box_reg: 0.0841 (0.1048) loss_mask: 0.2356 (0.2699) loss_objectness: 0.0177 (0.0241)
loss_rpn_box_reg: 0.0113 (0.0151) time: 0.4120 data: 0.0107 max mem: 6368
Epoch: [8] Total time: 0:00:25 (0.4263 s / it)
Epoch: [9] [ 0/60] eta: 0:00:53 lr: 0.0000005 loss: 0.7721 (0.7721) loss_classifier: 0.1039 (0.1039)
loss_box_reg: 0.2174 (0.2174) loss_mask: 0.3776 (0.3776) loss_objectness: 0.0391 (0.0391)
loss_rpn_box_reg: 0.0340 (0.0340) time: 0.8837 data: 0.4085 max mem: 6368
Epoch: [9] [10/60] eta: 0:00:22 lr: 0.0000005 loss: 0.4864 (0.5357) loss_classifier: 0.0601 (0.0711)
loss_box_reg: 0.1295 (0.1232) loss_mask: 0.2786 (0.2975) loss_objectness: 0.0300 (0.0271)
loss_rpn_box_reg: 0.0168 (0.0167) time: 0.4527 data: 0.0444 max mem: 6368
Epoch: [9] [20/60] eta: 0:00:17 lr: 0.0000005 loss: 0.4764 (0.5096) loss_classifier: 0.0678 (0.0685)
loss_box_reg: 0.1135 (0.1182) loss_mask: 0.2575 (0.2840) loss_objectness: 0.0215 (0.0234)
loss_rpn_box_reg: 0.0166 (0.0156) time: 0.4192 data: 0.0097 max mem: 6368
Epoch: [9] [30/60] eta: 0:00:12 lr: 0.0000005 loss: 0.4386 (0.4874) loss_classifier: 0.0575 (0.0649)
loss_box_reg: 0.0879 (0.1082) loss_mask: 0.2436 (0.2752) loss_objectness: 0.0169 (0.0249)
loss_rpn_box_reg: 0.0108 (0.0141) time: 0.4190 data: 0.0130 max mem: 6368
Epoch: [9] [40/60] eta: 0:00:08 lr: 0.0000005 loss: 0.4371 (0.4878) loss_classifier: 0.0575 (0.0661)
loss_box_reg: 0.0819 (0.1087) loss_mask: 0.2608 (0.2741) loss_objectness: 0.0168 (0.0243)
loss_rpn_box_reg: 0.0140 (0.0147) time: 0.4219 data: 0.0160 max mem: 6368
Epoch: [9] [50/60] eta: 0:00:04 lr: 0.0000005 loss: 0.4522 (0.4903) loss_classifier: 0.0597 (0.0657)
loss_box_reg: 0.0946 (0.1086) loss_mask: 0.2651 (0.2769) loss_objectness: 0.0174 (0.0241)
loss_rpn_box_reg: 0.0142 (0.0150) time: 0.4261 data: 0.0140 max mem: 6368
Epoch: [9] [59/60] eta: 0:00:00 lr: 0.0000005 loss: 0.4792 (0.4862) loss_classifier: 0.0574 (0.0654)

```

```
loss_box_reg: 0.0962 (0.1089) loss_mask: 0.2667 (0.2742) loss_objectness: 0.0156 (0.0228)
loss_rpn_box_reg: 0.0135 (0.0149) time: 0.4234 data: 0.0116 max mem: 6368
Epoch: [9] Total time: 0:00:25 (0.4304 s / it)
```

Evaluate the model (Option 2: Mobilenet)

```
evaluate(model_2, data_loader_test, device=device)
```

```
creating index...
```

```
index created!
```

```
Test: [ 0/50] eta: 0:00:18 model_time: 0.1628 (0.1628) evaluator_time: 0.0078 (0.0078)
      time: 0.3725 data: 0.2006 max mem: 6368
```

```
Test: [49/50] eta: 0:00:00 model_time: 0.0530 (0.0586) evaluator_time: 0.0085 (0.0116)
      time: 0.0707 data: 0.0049 max mem: 6368
```

```
Test: Total time: 0:00:04 (0.0826 s / it)
```

```
Averaged stats: model_time: 0.0530 (0.0586) evaluator_time: 0.0085 (0.0116)
```

```
Accumulating evaluation results...
```

```
DONE (t=0.01s).
```

```
Accumulating evaluation results...
```

```
DONE (t=0.01s).
```

```
IoU metric: bbox
```

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.441
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.907
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.307
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.130
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.462
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.234
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.551
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.559
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.375
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.572
```

```
IoU metric: segm
```

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.356
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.837
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.260
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.035
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.379
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.214
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.442
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.448
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.150
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.469
```

Now that training has finished, let's have a look at what it actually predicts in a test image

```
from PIL import Image
import requests
from torchvision import transforms

im = Image.open(requests.get(
    'https://upload.wikimedia.org/wikipedia/en/4/42/Beatles_-_Abbey_Road.jpg',
    stream=True).raw)
convert_tensor = transforms.ToTensor()
x = convert_tensor(im)

Image.fromarray(x.mul(255).permute(1, 2, 0).byte().numpy())
```



And let's now visualize the top predicted segmentation mask. The masks are predicted as $[N, 1, H, W]$, where N is the number of predictions, and are probability maps between 0-1.

```
model_1.eval()
model_2.eval()
with torch.no_grad():
    prediction_1 = model_1([x.to(device)])
    prediction_2 = model_2([x.to(device)])
```

Model 1 Output Visualization (Top 3 Masks).

```
im1 = Image.fromarray(prediction_1[0]['masks'][0, 0].mul(255).byte().cpu().numpy())
im2 = Image.fromarray(prediction_1[0]['masks'][1, 0].mul(255).byte().cpu().numpy())
im3 = Image.fromarray(prediction_1[0]['masks'][2, 0].mul(255).byte().cpu().numpy())

im1.show()
im2.show()
im3.show()
```








Model 2 Output Visualization (Top 3 Masks).

```
im1 = Image.fromarray(prediction_2[0]['masks'][0, 0].mul(255).byte().cpu().numpy())
im2 = Image.fromarray(prediction_2[0]['masks'][1, 0].mul(255).byte().cpu().numpy())
im3 = Image.fromarray(prediction_2[0]['masks'][2, 0].mul(255).byte().cpu().numpy())

im1.show()
im2.show()
im3.show()
```





(b)

Performance comparison between the two object detection models based on the provided evaluation metrics:

Average Precision (AP): AP measures the accuracy of the model in localizing the objects in the image. It is the area under the precision-recall curve.

Model 1 (ResNet) outperforms Model 2 (MobileNet) in terms of AP. The ResNet model has higher AP values across all the IoU thresholds and all object sizes (small, medium, and large). The ResNet model achieves an AP of 0.990 at IoU=0.50 for all object sizes, whereas the MobileNet model achieves an AP of 0.907. Similarly, the ResNet model achieves an AP of 0.846 at IoU=0.50:0.95 for large objects, whereas the MobileNet model achieves an AP of 0.462.

Average Recall (AR): AR measures the ability of the model to detect all instances of the object in the image. It is the area under the recall-IoU curve.

Model 1 (ResNet) outperforms Model 2 (MobileNet) in terms of AR. The ResNet model has higher AR values across all the IoU thresholds and all object sizes (small, medium, and large). The ResNet model achieves an AR of 0.876 at IoU=0.50:0.95 for all object sizes, whereas the MobileNet model achieves an AR of 0.802. Similarly, the ResNet model achieves an AR of 0.883 at IoU=0.50:0.95 for large objects, whereas the MobileNet model achieves an AR of 0.807.

Intersection over Union (IoU): IoU measures the overlap between the predicted bounding box and the ground truth bounding box. It is the ratio of the intersection area to the union area.

Model 1 (ResNet) outperforms Model 2 (MobileNet) in terms of IoU. The ResNet model has higher IoU values across all the object sizes (small, medium, and large) for both bbox and segm metrics. For example, the ResNet model achieves an IoU of 0.960 at IoU=0.75 for all object sizes for bbox metric, whereas the MobileNet model achieves an IoU of 0.307. Similarly, the ResNet model achieves an IoU of 0.918 at IoU=0.75 for all object sizes for segm metric, whereas the MobileNet model achieves an IoU of 0.130.

In summary, based on the evaluation metrics provided, the ResNet model outperforms the MobileNet model in terms of object detection accuracy. The ResNet model achieves higher values of AP, AR, and IoU across all object sizes and IoU thresholds.

(c)

We tested two different models for image segmentation, referred to as "model 1" and "model 2". Model 1 used a more complex architecture called "ResNet" while model 2 used a simpler architecture called "MobileNet".

When testing the models using a real image consisting of humans to be segmented, they found that the segmentation produced by model 1 was sharper, while the segmentation produced by model 2 was smoother. This difference in performance could be attributed to the fact that ResNet is a more complex architecture with more layers and parameters, allowing it to capture more intricate details in the image. On the other hand, MobileNet is a simpler architecture designed for mobile devices, which may not be as powerful or accurate as DeepLabV3+ in certain tasks.

Overall, the findings suggest that the choice of architecture can have a significant impact on the performance of image segmentation models. Depending on the specific task and requirements, a more complex or simpler architecture may be more appropriate.